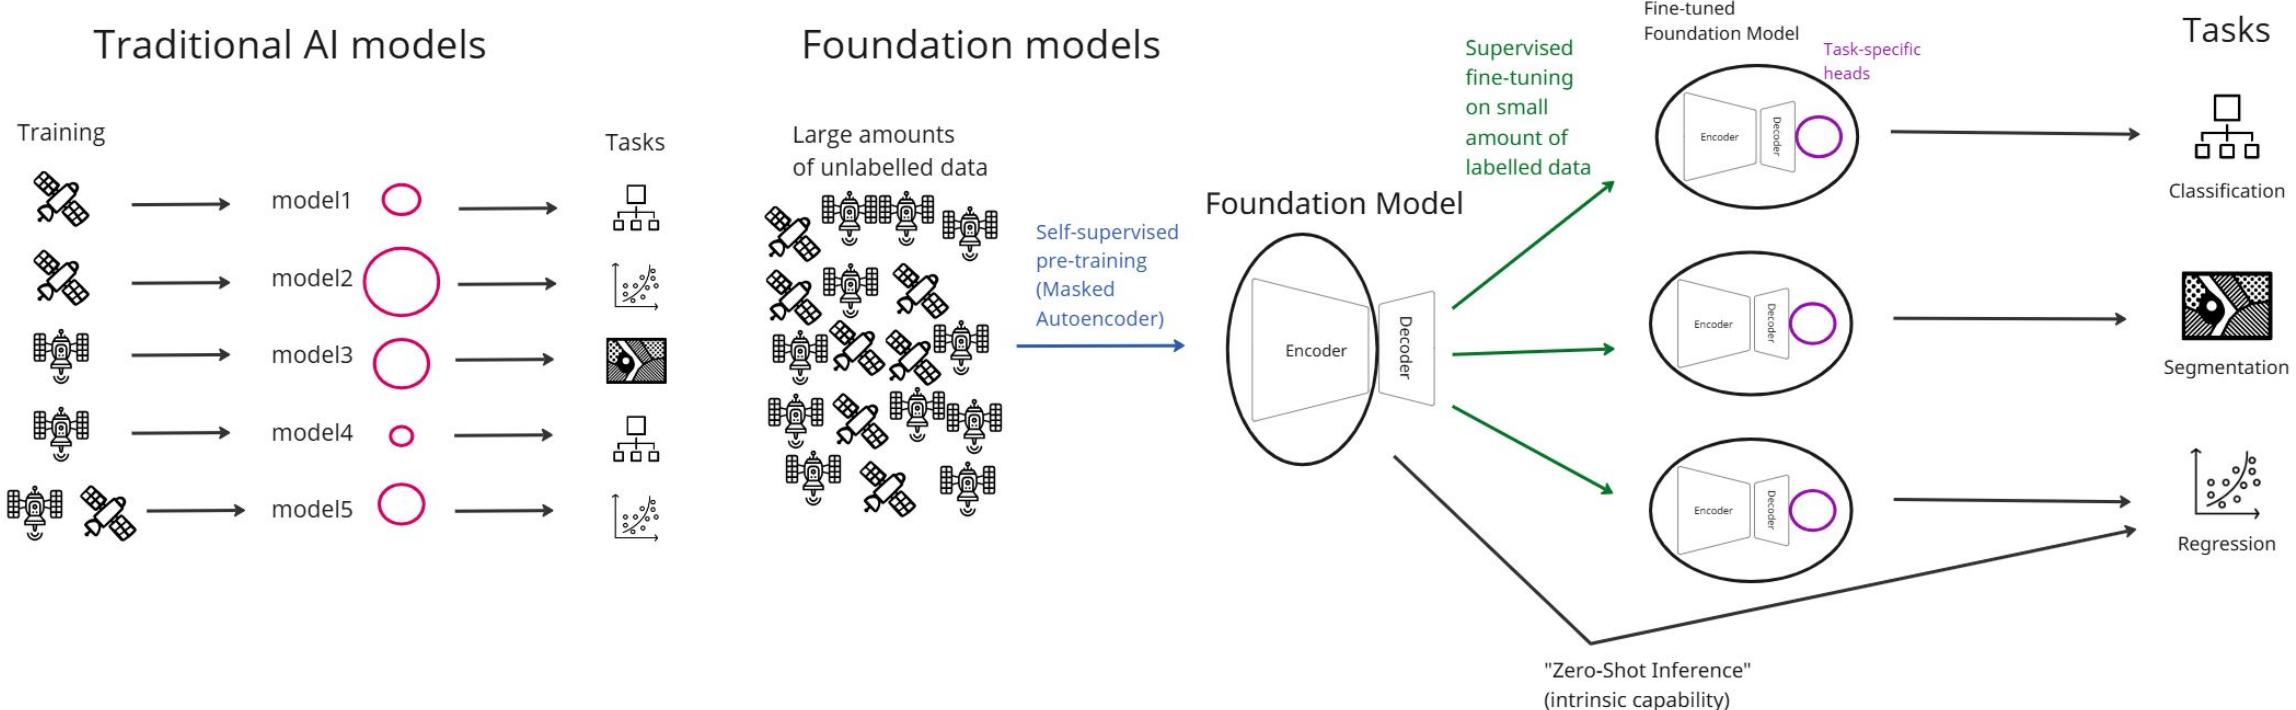


# Cloud Masking Sen-3 with Prithvi-EO Foundation Model

Boran Frank  
*Intern in Copernicus Programme Office*

23.04.2025





## Advantages of using FMs:

- Harness large amounts of unlabelled data
- Fine-tuning requires only a small amount of labelled data
- Adaptation to different tasks more efficient than building new standalone models from scratch
- Supported by Python frameworks like **TerraTorch**

**Extended Data Table 1:** State-of-the-art EO FMs. The slash sign / indicates separate models for each; the range sign - indicates a rough range as the exact numbers are not available; the comma sign , indicates a combined coverage.

Model	Pre-training Dataset	Data Size (# pixels)	Modality	Spatial Resolution	Spatial Coverage (km <sup>2</sup> )	Model Size (# params)	Temporal Coverage	Sequence Length
SeCo <sup>33</sup>	SeCo	70B	MS	10m	Global (1.4M)	23M	N/A	1
SatMAE <sup>46</sup>	fMoW	50B	MS / RGB	0.5m / 10m	Global (1.7M)	305M	5 years	3
RVSA <sup>122</sup>	MillionAID	4B	RGB	0.5m-150m	Global (18K)	89M	N/A	1
Scale-MAE <sup>45</sup>	fMoW-RGB	80B	RGB	0.3-10m	Global (1.7M)	305M	N/A	1
CROMA <sup>49</sup>	SSL4EO-S12	140B	MS, SAR	10m	Global (1.8M)	305M	N/A	1
DeCUR <sup>50</sup>	SSL4EO-S12 / GeoNRW	140B	MS, SAR / RGB, DEM	10m / 1m	Global (1.8M)	23M	N/A	1
Presto <sup>47</sup>	Presto	4B	MS, SAR, ERA5, DEM, etc.	10m	Global (2K)	400K	2 years	24
Prithvi <sup>48</sup>	HLS	158B	MS	10m, 30m	US	100M	1+ years	3
SpectralGPT <sup>123</sup>	fMoW+BigEarthNet	12B	MS	10m	Global (1.7M)	86M	N/A	1
SkySense <sup>41</sup>	SkySense	97T	MS, RGB, SAR	0.3m, 10m	Global (8.8M)	2B	?	20
OFA-Net <sup>51</sup>	Satlas.+5B pix. +HypspecNet.	13B	MS, RGB, SAR, HS	1m, 4m, 10m, 30m	Global (260K)	86M	N/A	1
DOFA <sup>52</sup>	Satlas.+5B pix. +HypspecNet.	105B	MS, RGB, SAR, HS	1m, 4m, 10m, 30m	Global (2.6M)	305M	N/A	1
SSL4EO-S12 <sup>34</sup>	SSL4EO-S12	140B	MS, SAR	10m	Global (1.8M)	23M	N/A	1
SSL4EO-L <sup>35</sup>	SSL4EO-L	348B	MS	30m	Global (15.6M)	23M	N/A	1
SatlasPretrain <sup>36</sup>	SatlasPretrain	17T	MS, RGB, SAR	0.5-2m, 10m	Global (21M)	88M	11 years	10
SkyScript CLIP <sup>54</sup>	SkyScript	130B	MS, RGB, text	0.5m, 1-2m, 10-30m	Global (5M)	305M	N/A	1

**Extended Data Table 2:** State-of-the-art weather and climate FMs.

Model	Affiliation	Model Architecture	Temporal Resolution	Time Span	Spatial Resolution	Spatial Coverage	Model Size (# params)	Cluster Size	Training Time
KeislerNet <sup>124</sup>	Descartes	GNN	6 hrs	6 days	1°	Global	6.7M	1 A00	5.5 days
FourCastNet <sup>125</sup>	NVIDIA	Transformer/FNO	6 hrs	200 hrs	0.25°	Global	433M-2.19B	64 A100	16 hrs
Pangu-Weather <sup>67</sup>	Huawei	Transformer	1 hr	7 days	0.25°	Global	256M	192 V100	15 days
GraphCast <sup>68</sup>	Google	GNN	6 hrs	10 days	0.25°	Global	36.7M	32 TPU v4	4 wks
ClimaX <sup>63</sup>	UCLA	Transformer	6 hrs	1 month	0.25°	Global	108-111M	80 V100	3 days
FengWu <sup>126</sup>	Shanghai	Transformer	6 hrs	14 days	0.25°	Global	427M	32 A100	17 days
W-MAE <sup>127</sup>	Chengdu	Transformer	6 hrs	9 days	0.25°	Global	96.5M	8 A800	223.4 hrs
FuXi <sup>91</sup>	Fudan	Transformer	6 hrs	15 days	0.25°	Global	1.56B	8 A100	30 hrs
SEEDS <sup>128</sup>	Google	Diffusion	N/A	16 days	2°	Global	114M	16 TPU v4	18 hrs
NeuralGCM <sup>69</sup>	Google	Differential GCM	3.75-12 mins	15 days	0.7-2.8°	Global	115-311M	128 TPU v5e	10 days
Stormer <sup>129</sup>	UCLA	Transformer	6 hrs	14 days	1.4°	Global	100-600M	128 A100	24 hrs
GenCast <sup>92</sup>	Google	Diffusion	12 hrs	15 days	1°	Global	60M	32 TPU v4	5 days
FengWu-GHR <sup>59</sup>	Shanghai	Transformer	6 hrs	10 days	0.09°	Global	4B	16 A100	3 days

Zhu et al. 2024

□ What about medium-/low-resolution Earth Observation (e.g. Sen-3, Meteosat ...)?

## Earth Observation Foundation Models

- trained on high-resolution data (e.g. Sentinel-2)
- primarily land applications

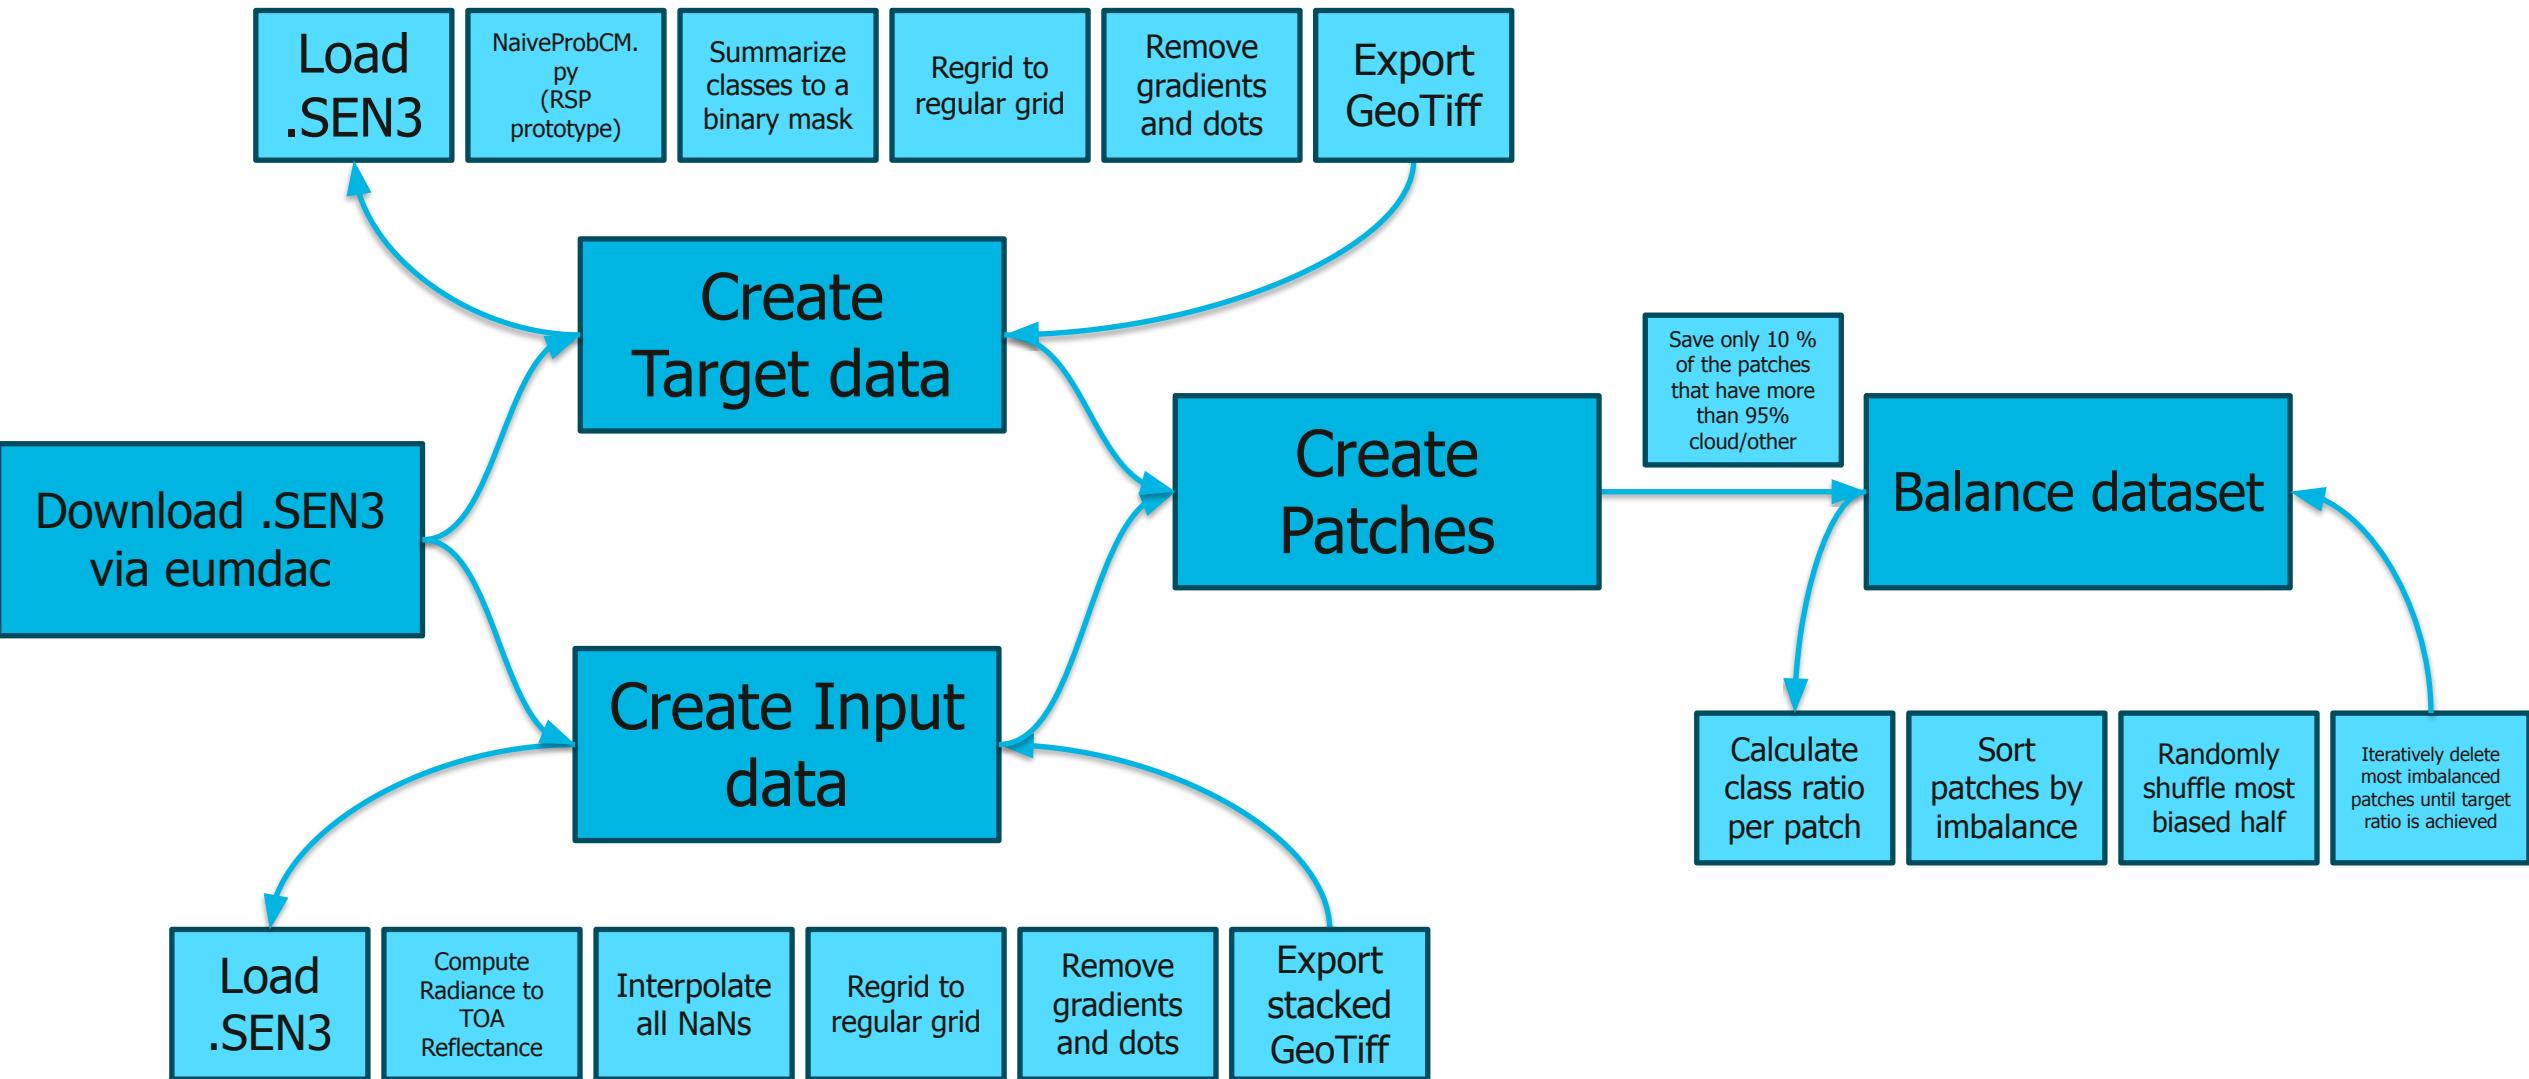
## Weather and Climate Foundation Models

- trained on low-resolution reanalysis data

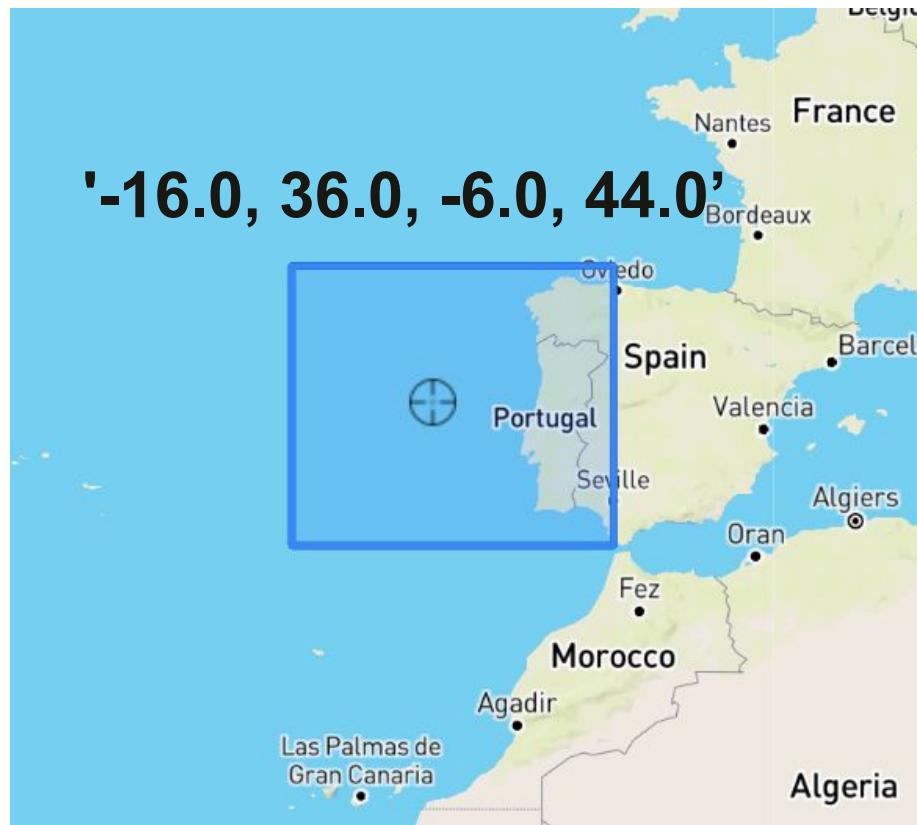


## Research Questions:

- How does a foundation model perform that was pre-trained with a higher resolution?
  - How does a foundation model perform that was pre-trained on other spectral bands?
  - How easy is it to apply foundation models? What computational and intellectual resources are required?
- Do this with the Prithvi-EO 2.0 foundation model



Randomly downloading tiles from summer 2024, 2023 that overlap this bounding box





# Methodology: Choose bands

S2 band number used in Prithvi-EO	S2 corresponding wavelength	S3 SLSTR band number	S3 SLSTR corresponding wavelength	Used in Prithvi_CM
2	0.49 (Blue)	None, but OLCI 4	None, but OLCI4: 0.49	2 (0.66)
3	0.56 (Green)	1	0.55	2 (0.66)
4	0.665 (Red)	2	0.66	3 (0.87)
8A	0.865 (NIR)	3	0.87	5 (1.61)
11	1.610 (SWIR1)	5	1.61	5 (1.61)
12	2.190 (SWIR2)	6	2.25	6 (2.25)

**BIG  
OOPS!**

Extended research question:

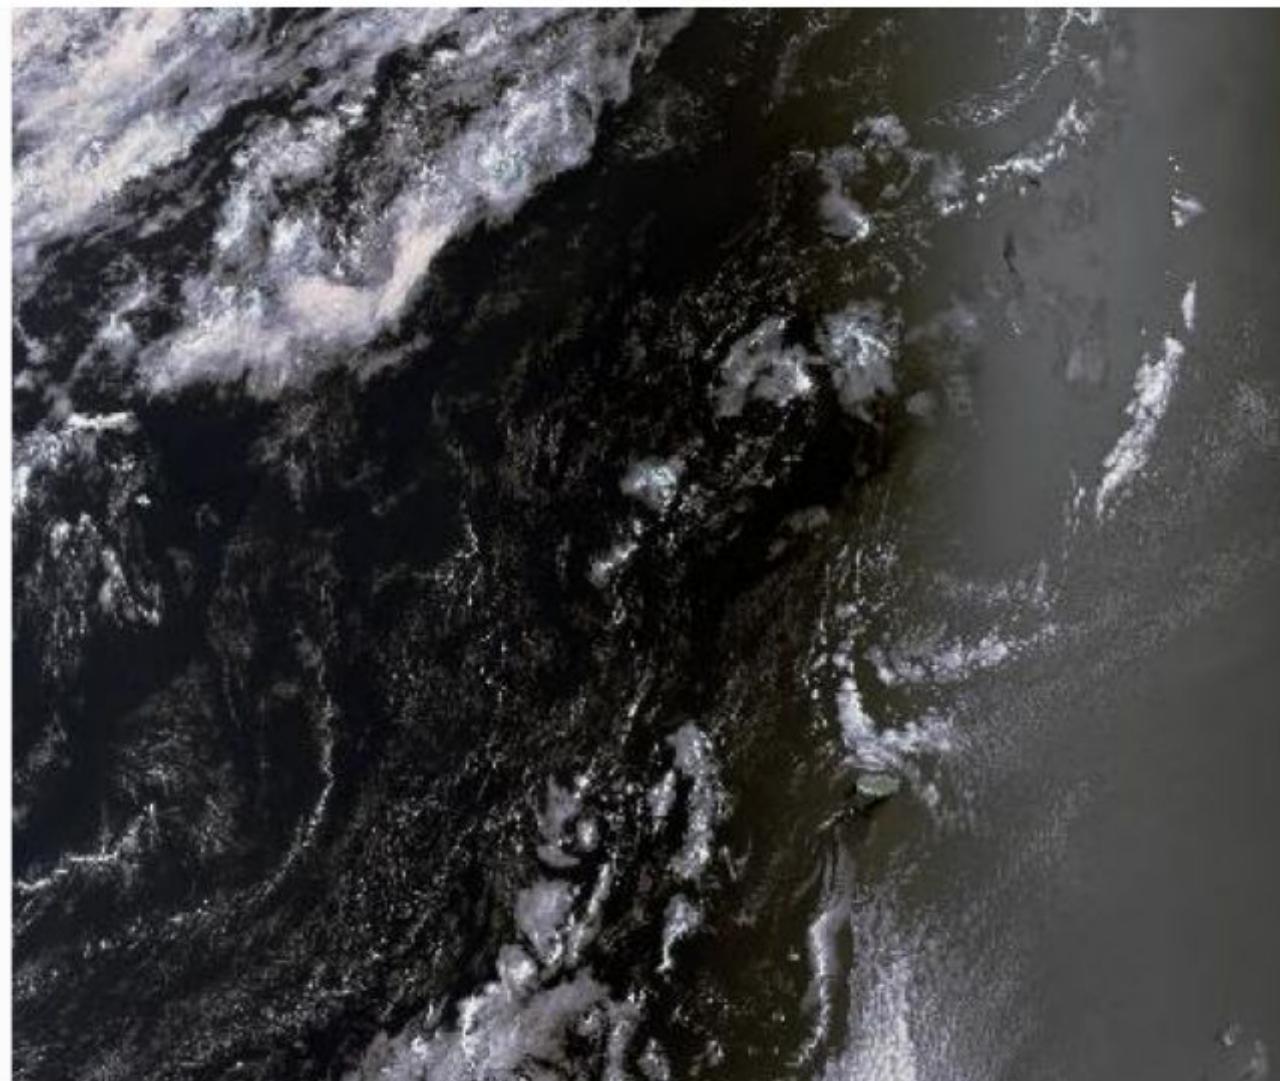
How does the FM perform with suboptimally chosen spectral bands?



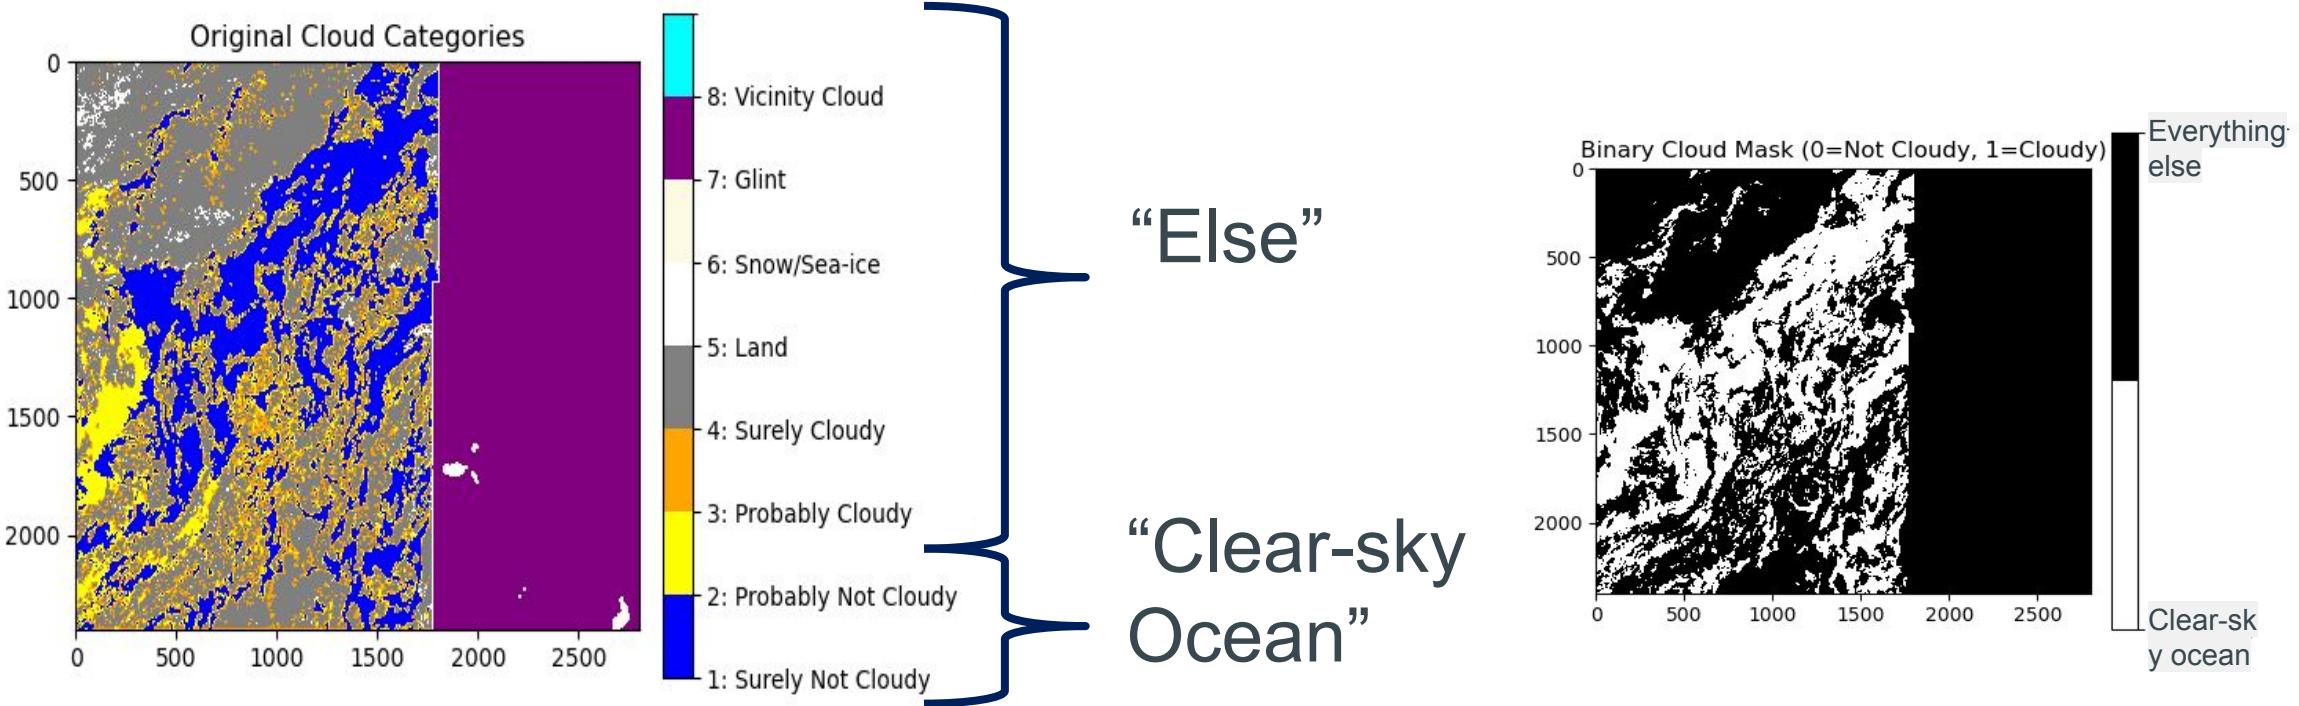
# Methodology: Example data: 20230721T104815

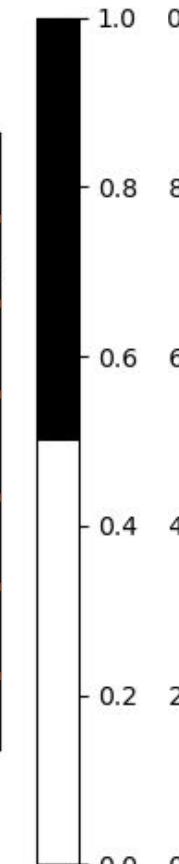
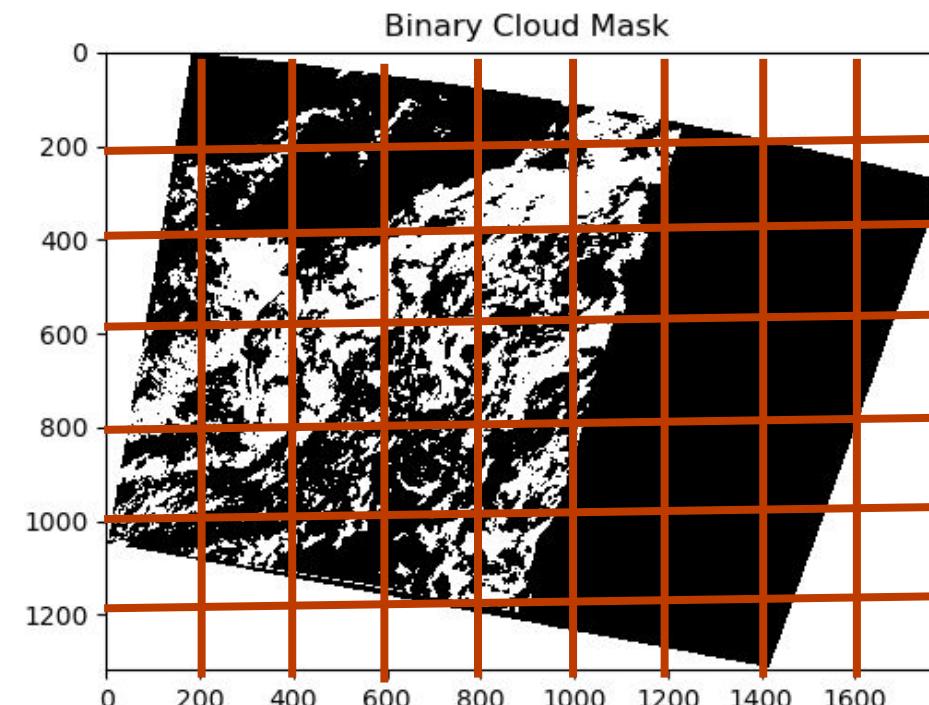
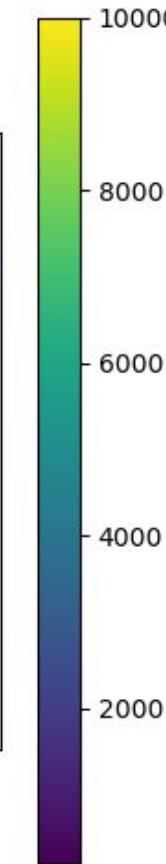
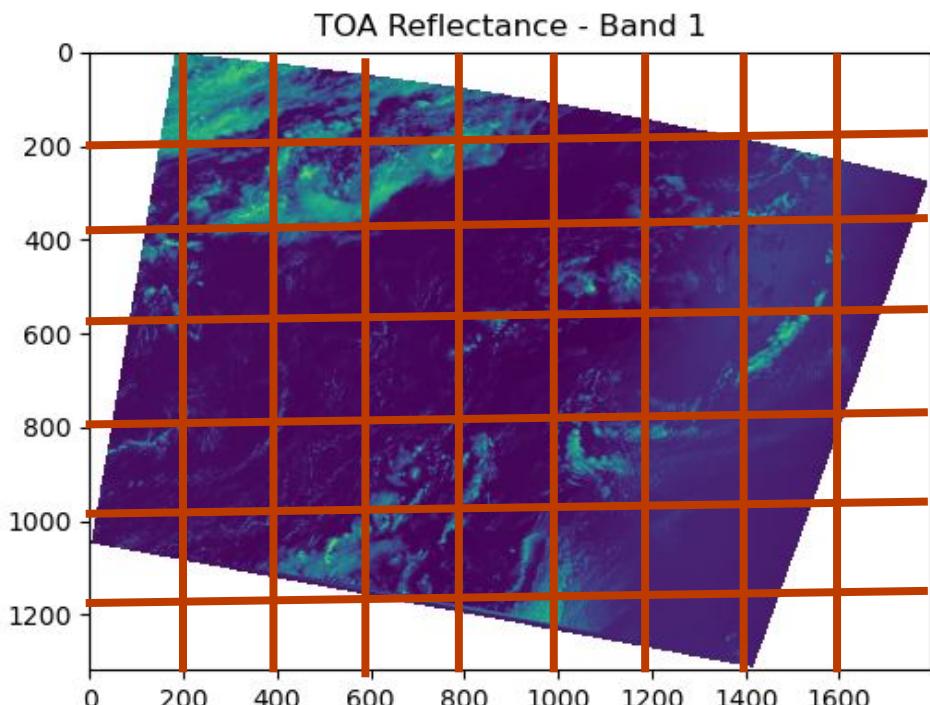
[www.eumetsat.int](http://www.eumetsat.int)

Browse Image



Nadir view!

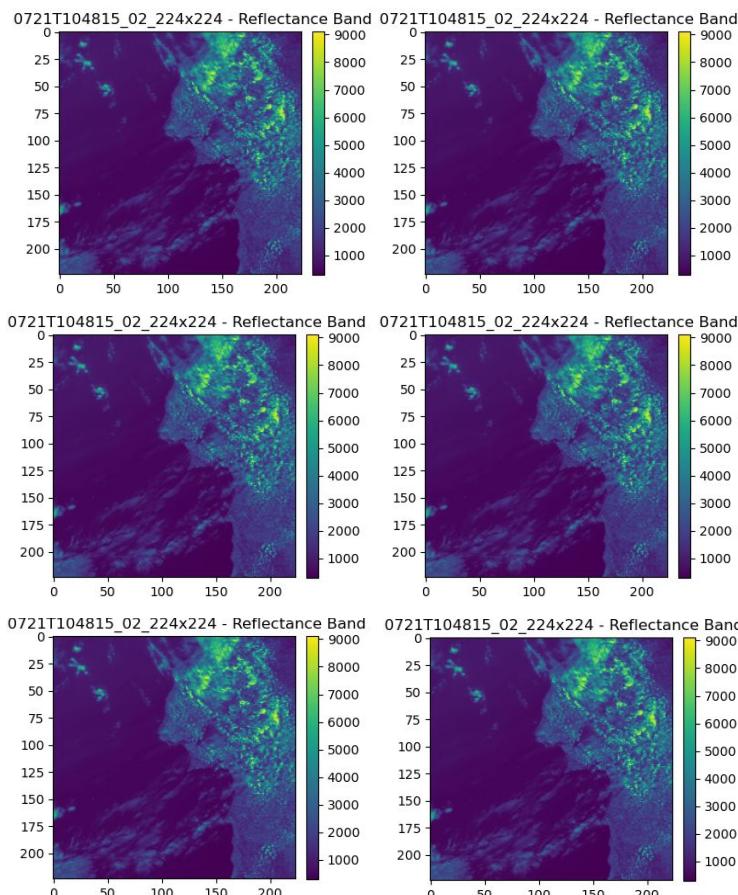




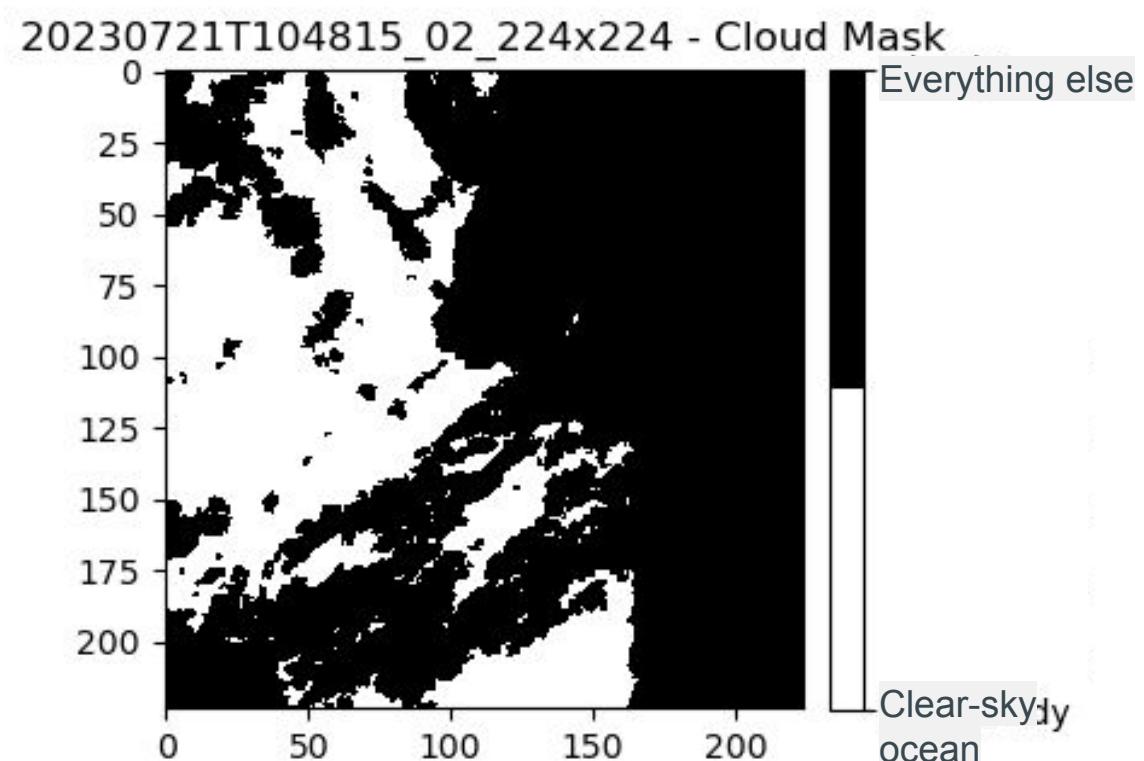
224 x 224 patches are cut out

re-gridding introduces NAN borders  only patches are used that do not include any NAN border  not optimal

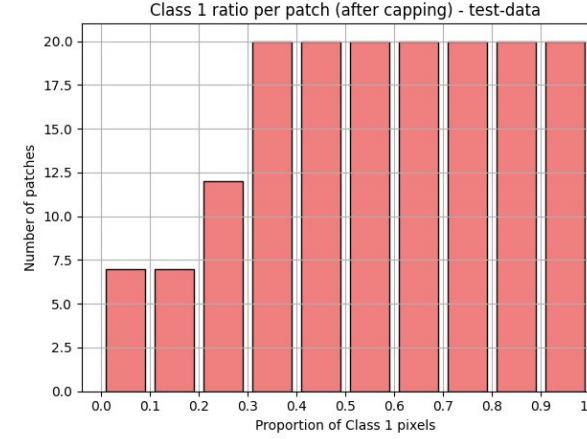
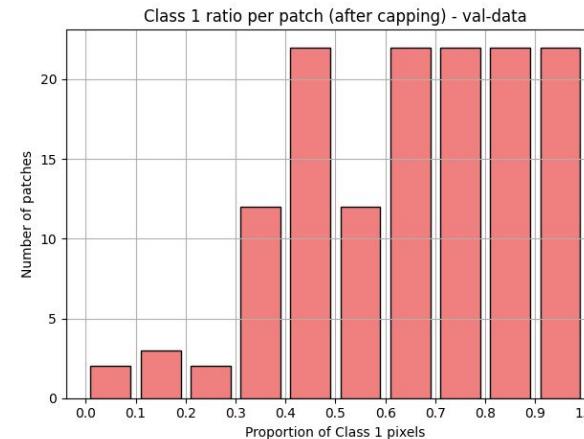
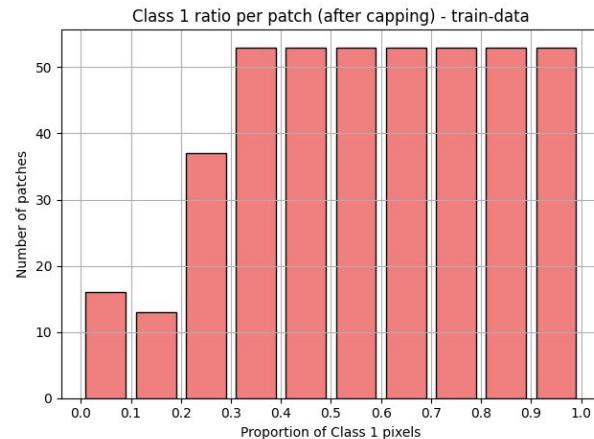
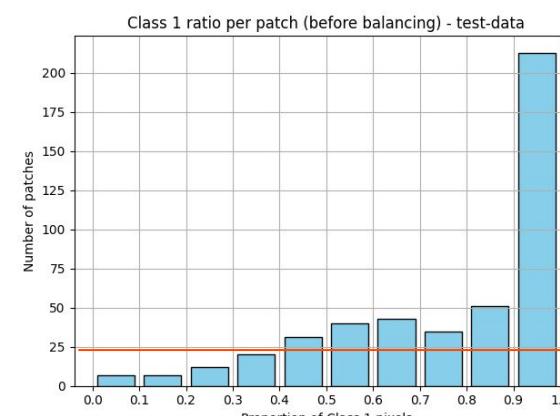
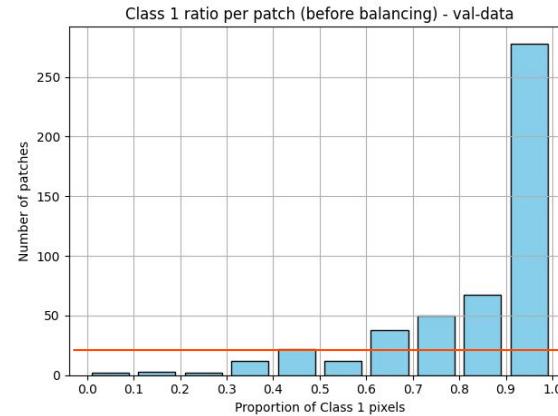
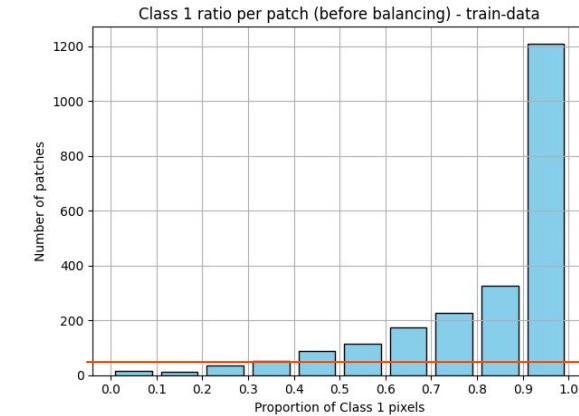
INPUT: SLSTR patches from bands: 2,2,3,5,5,6



TARGET: Binary Cloud Mask from  
NaiveProbCM.py



! Here only band 1 is displayed, should be 2,2,3,5,5,6



## Final dataset

Train: 437

Val: 141

Test: 166

### Class distribution in 'train':

Class 0: 9096526 pixels  
(41.49%)

Class 1: 12830386 pixels  
(58.51%)

### Class distribution in 'val':

Class 0: 2384375 pixels  
(33.70%)

Class 1: 4690441 pixels  
(66.30%)

### Class distribution in 'test':

Class 0: 3489298 pixels  
(41.89%)

Class 1: 4839918 pixels  
(58.11%)



## Prithvi-EO:

- Developed by IBM x NASA
- trained on 4.2m 256 x 256 samples of HLS (Harmonized Landsat Sentinel-2) with **30m** resolution
- Vision Transformer with 300M or 600M parameters
  - Trained with Masked Autoencoder strategy
- trained for ~ 58k GPU hours (NVIDIA A100)
- **[2412.02732] Prithvi-EO-2.0: A Versatile Multi-Temporal Foundation Model for Earth Observation Applications**

## TerraTorch:

- flexible fine-tuning framework for Geospatial Foundation Models based on **TorchGeo** ([torchgeo](#)) and **PyTorch Lightning** ([PyTorch Lightning](#))
- open-source code and weights
- easily access pre-trained and fine-tuned models, datasets, configurations
- **[IBM/terratorch: A Python toolkit for fine-tuning Geospatial Foundation Models \(GFMs\).](#)**



# Methodology: Workflow with TerraTorch

## Define datamodule

```
datamodule = terratorch.datamodules.GenericNonGeoSegmentationDataModule(  
    batch_size=8,  
    num_workers=2,  
    num_classes=2,  
  
    # Define dataset paths  
    train_data_root=dataset_path / 'data/',  
    train_label_data_root=dataset_path / 'data/',  
    val_data_root=dataset_path / 'data/',  
    val_label_data_root=dataset_path / 'data/',  
    test_data_root=dataset_path / 'data/',  
    test_label_data_root=dataset_path / 'data/',  
  
    # Define splits  
    train_split=dataset_path / 'splits/train.txt',  
    val_split=dataset_path / 'splits/val.txt',  
    test_split=dataset_path / 'splits/test.txt',  
  
    img_grep='*_reflectance.tif',  
    label_grep='*_binary.tif',  
  
    train_transform=[  
        albumentations.D4(), # Random flips and rotation  
        albumentations.pytorch.transforms.ToTensorV2(),  
    ],  
    val_transform=None, # Using ToTensor() by default  
    test_transform=None,  
  
    # Define standardization values  
    means=[  
        stats[0][0],  
        stats[1][0],  
        stats[2][0],  
        stats[3][0],  
        stats[4][0],  
        stats[5][0],  
    ],  
    stds=[  
        stats[0][1],  
        stats[1][1],  
        stats[2][1],  
        stats[3][1],  
        stats[4][1],  
        stats[5][1],  
    ],  
    no_data_replace = 0,  
    no_label_replace = -1,  
    # We use all six bands of the data, so we don't need to define dataset_bands  
)
```

## Define model

```
# Model  
model = terratorch.tasks.SemanticSegmentationTask(  
    model_factory="EncoderDecoderFactory",  
    model_args={  
        # Backbone  
        "backbone": "prithvi_eo_v2_300_tl", # Model can be either prithvi_eo_v1_100, prithvi_eo_v2_300  
        "backbone_pretrained": True,  
        "backbone_num_frames": 1, # 1 is the default value  
        "backbone_bands": ["BLUE", "GREEN", "RED", "NIR_NARROW", "SWIR_1", "SWIR_2"],  
        "backbone_coords_encoding": ['time', 'location'], # use ["time", "location"] for time and location  
  
        # Necks  
        "necks": [  
            {  
                "name": "SelectIndices",  
                "# indices": [2, 5, 8, 11] # indices for prithvi_eo_v1_100  
                "indices": [5, 11, 17, 23] # indices for prithvi_eo_v2_300  
                "# indices": [7, 15, 23, 31] # indices for prithvi_eo_v2_600  
            },  
            {"name": "ReshapeTokensToImage"},  
            {"name": "LearnedInterpolateToPyramidal"}  
        ],  
  
        # Decoder  
        "decoder": "UNetDecoder",  
        "decoder_channels": [512, 256, 128, 64],  
  
        # Head  
        "head_dropout": 0.1,  
        "num_classes": 2,  
    },  
  
    loss="dice",  
    optimizer="AdamW",  
    lr=1e-4,  
    ignore_index=-1,  
    freeze_backbone=True, # Only to speed up fine-tuning  
    freeze_decoder=False,  
    plot_on_val=True,  
    class_names=['no cloud', 'cloud o.o.']) # optionally define class names
```

## Define trainer

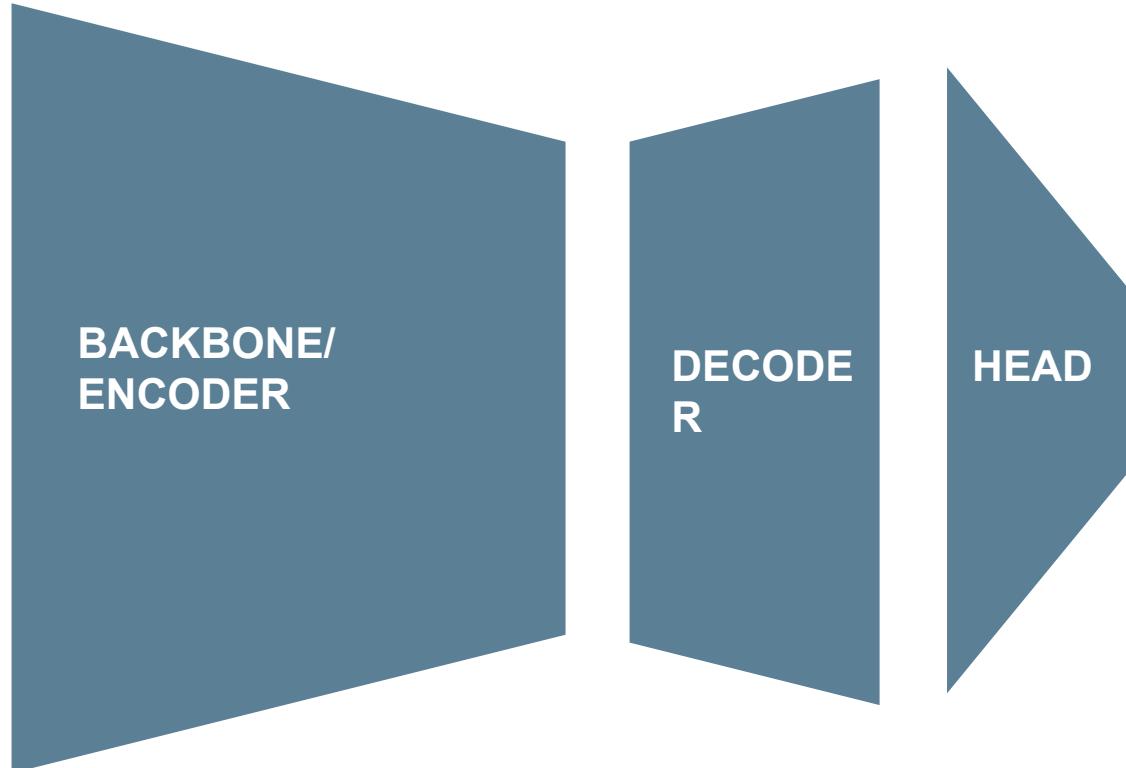
```
# Lightning Trainer  
trainer = pl.Trainer(  
    accelerator="auto",  
    strategy="auto",  
    devices=1, # Deactivate multi-gpu because it often fails in notebooks  
    precision='16-mixed', # Speed up training  
    num_nodes=1,  
    logger=True, # Uses TensorBoard by default  
    max_epochs=15,  
    log_every_n_steps=1,  
    enable_checkpointing=True,  
    callbacks=[checkpoint_callback, pl.callbacks.RichProgressBar()],  
    default_root_dir="output/sen3_cm",  
)
```



# Methodology: Model Architecture and Training strategy

www.eumetsat.int

- Prithvi backbone (Vision Transformer):  
303M or 631M parameters
- U-Net decoder: 20.3M or 25.7M
- Segmentation head (ConvNN): 4.8M
- Total size: 1.3 GB or 2.6 GB
- Epochs = 20
- Batch size = 8
- Lr = 1e-4
- AdamW
- Loss = dice
- vRAM needed: 8025MB for Prithvi\_300  
when all parameters are trained
- Time per epoch: ~30s



Pretrained\_encoder = True  Foundation model/ pre-trained weights are used

Pretrained\_encoder = False  Foundation model/ pre-trained weights are NOT used  
 randomly initialized

Frozen\_encoder = True  Encoder weights are NOT changed during fine-tuning

Frozen\_encoder = False  Encoder weights are changed during fine-tuning

Frozen\_decoder = True  Decoder weights are NOT changed during fine-tuning

Frozen\_decoder = False  Decoder weights are changed during fine-tuning

! Decoder and Head always start with random weights

Dice:

- 0 = no overlap
- 1 = perfect overlap
- Dice loss = 1 / Dice Coefficient

$$\text{Dice Coefficient} = \frac{2 \times \text{Intersection}}{\text{Size of Predicted} + \text{Size of Actual}}$$

Multiclass Jaccard Index:

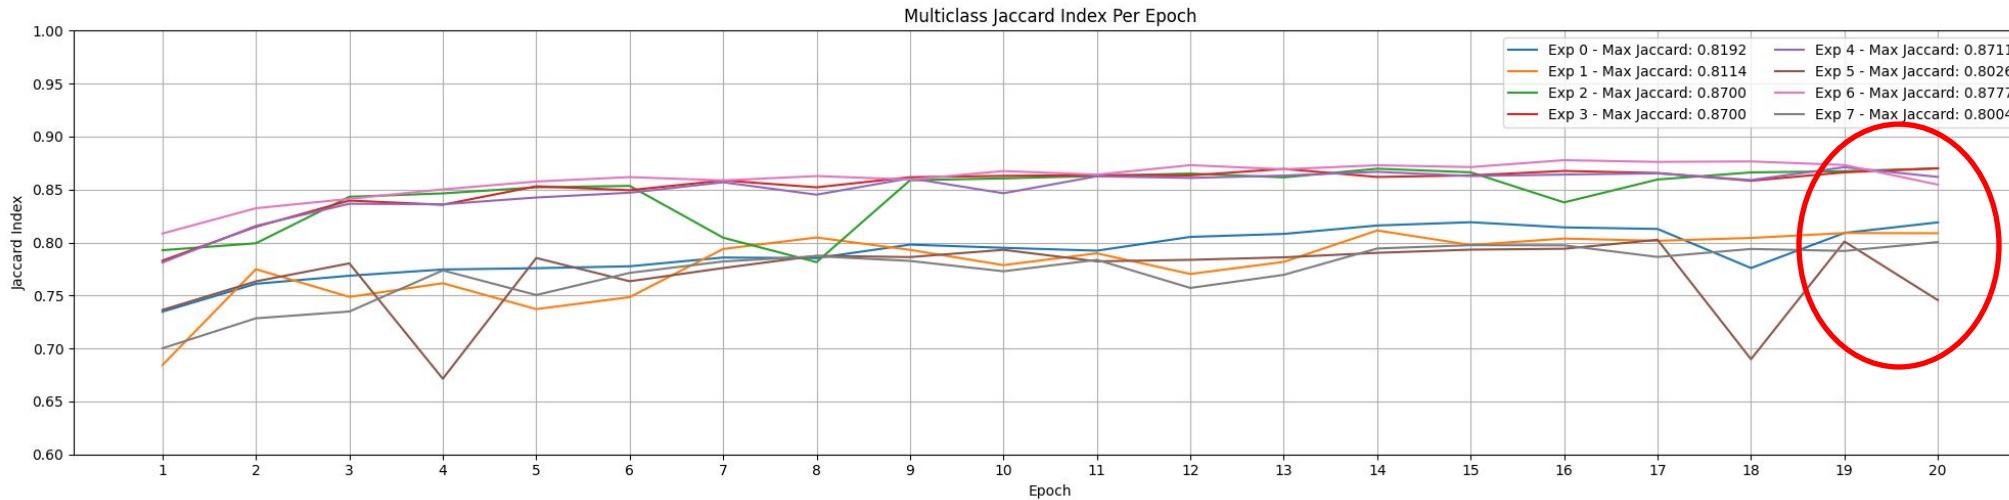
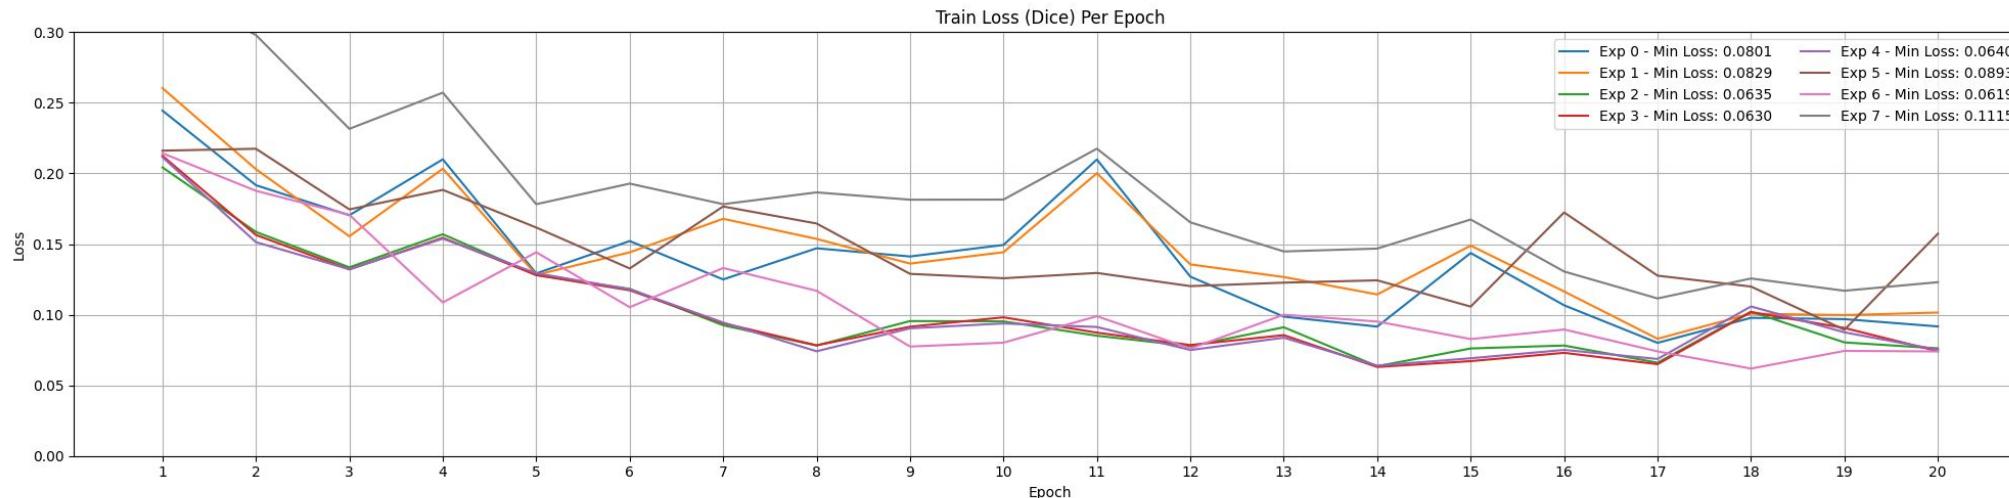
- 0 = no overlap
- 1 = perfect overlap

$$\text{Jaccard Index} = \frac{\text{Intersection of predicted and ground truth}}{\text{Union of predicted and ground truth}}$$



# Results: Ablation studies

experiment	backbone	backbone_pretrained	backbone_encoding	freeze_bac_kbone	freeze_decoder	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_0	v2_300	FALSE	[]	FALSE	FALSE	0.1088 (6.)	0.8226 (6.)
exp_1	v2_300	FALSE	[]	TRUE	FALSE	0.1191 (7.)	0.8103 (7.)
exp_2	v2_300	TRUE	[]	FALSE	FALSE	0.0774 (3.)	0.8697 (3.)
exp_3	v2_300	TRUE	[]	TRUE	FALSE	0.0901 (4.)	0.8576 (4.)
exp_4	v2_300_tl	TRUE	['time', 'location']	TRUE	FALSE	0.0725 (2.)	0.8798 (2.)
exp_5	v2_600	FALSE	[]	FALSE	FALSE	0.1001 (5.)	0.8313 (5.)
exp_6	v2_600	TRUE	[]	TRUE	FALSE	<b>0. 0711 (1.)</b>	<b>0.8836 (1.)</b>
exp_7	v2_300	FALSE	[]	TRUE	TRUE	0.1386 (8.)	0.8041 (8.)



Two groups:  
0,1,5,7, - worse □  
without pre-trained  
backbone  
2,3,4,6 – better □  
with pre-trained  
backbone

These comparisons are the key results, as they evaluate whether using pre-trained weights—essentially the foundation model—adds value.

experiment	backbone	backbone_pr etrained	backbone_enc oding	freeze_bac kbone	freeze_dec oder	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_0	v2_300	FALSE	False	False	False	0.1088 (6.)	0.8226 (6.)
exp_2	v2_300	TRUE	False	False	False	0.0774 (3.)	0.8697 (3.)
exp_1	v2_300	FALSE	False	True	False	0.1191 (7.)	0.8103 (7.)
exp_3	v2_300	TRUE	False	True	False	0.0901 (4.)	0.8576 (4.)

↑ 5.7 % improvement

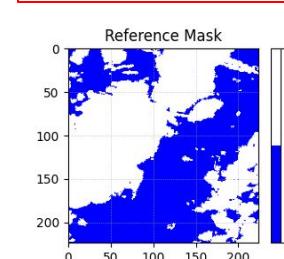
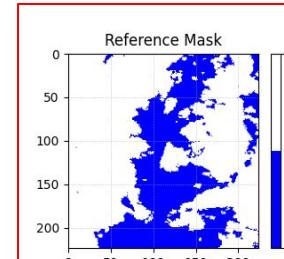
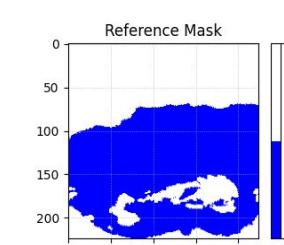
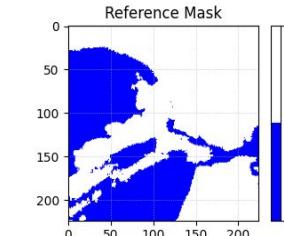
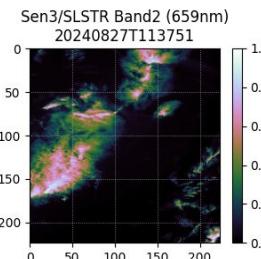
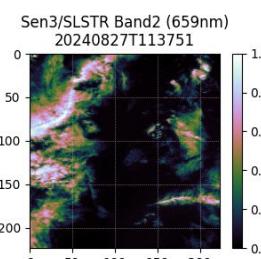
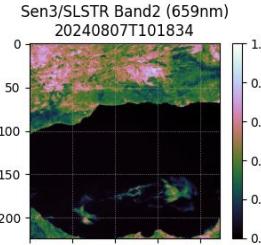
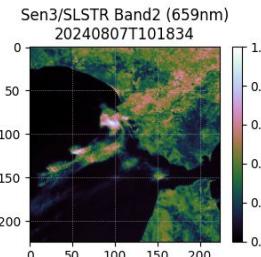
↑ 5.8 % improvement

# Main Results: Pre-trained backbone

0,1 : without  
pre-trained  
backbone

2,3: with  
pre-trained  
backbone

2,3 more  
detailed and  
recognizes  
small structures  
better



Exp: 0

3

Prediction  
Dice Loss: 0.030, Jaccard: 0.942

Prediction  
Dice Loss: 0.040, Jaccard: 0.923

Prediction  
Dice Loss: 0.050, Jaccard: 0.905

Prediction  
Dice Loss: 0.085, Jaccard: 0.843

2

Prediction  
Dice Loss: 0.020, Jaccard: 0.961

Prediction  
Dice Loss: 0.018, Jaccard: 0.966

Prediction  
Dice Loss: 0.032, Jaccard: 0.938

Prediction  
Dice Loss: 0.043, Jaccard: 0.917

1

Prediction  
Dice Loss: 0.030, Jaccard: 0.942

Prediction  
Dice Loss: 0.036, Jaccard: 0.930

Prediction  
Dice Loss: 0.048, Jaccard: 0.909

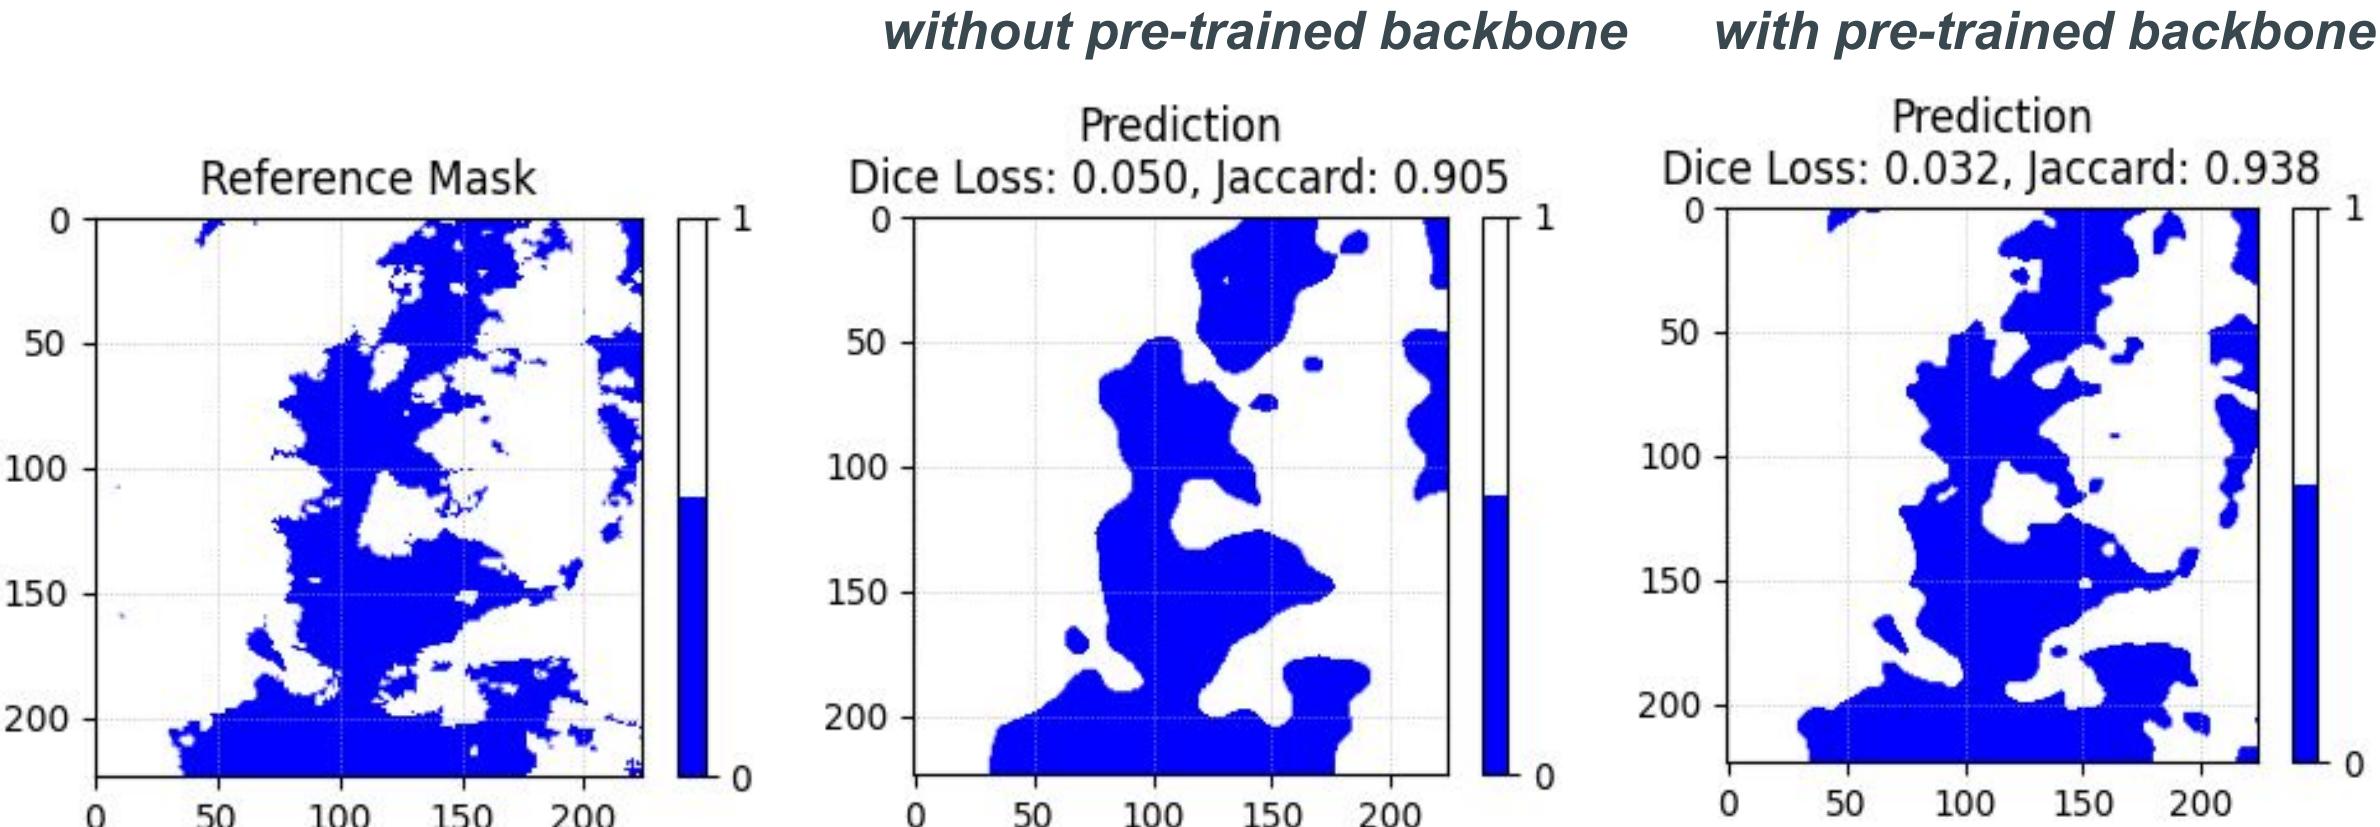
Prediction  
Dice Loss: 0.072, Jaccard: 0.866

Prediction  
Dice Loss: 0.017, Jaccard: 0.966

Prediction  
Dice Loss: 0.016, Jaccard: 0.968

Prediction  
Dice Loss: 0.029, Jaccard: 0.944

Prediction  
Dice Loss: 0.049, Jaccard: 0.907



experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_0	v2_300	FALSE	[]	FALSE	FALSE	0.1088 (6.)	0.8226 (6.)
exp_1	v2_300	FALSE	[]	TRUE	FALSE	0.1191 (7.)	0.8103 (7.)

- Non-frozen backbone without pre-trained weights improves MJI by 1,5 %

experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_2	v2_300	TRUE	[]	FALSE	FALSE	0.0774 (3.)	0.8697 (3.)
exp_3	v2_300	TRUE	[]	TRUE	FALSE	0.0901 (4.)	0.8576 (4.)

- Non-frozen backbone with pre-trained weights improves MJI by 1.4 %

experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_3	v2_300	TRUE	[]	TRUE	FALSE	0.0901 (4.)	0.8576 (4.)
exp_4	v2_300_tl	TRUE	['time', 'location']	TRUE	FALSE	0.0725 (2.)	0.8798 (2.)

- Using time and location embedding improves MJI by 2,5 %

experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_0	v2_300	FALSE	[]	FALSE	FALSE	0.1088 (6.)	0.8226 (6.)
exp_5	v2_600	FALSE	[]	FALSE	FALSE	0.1001 (5.)	0.8313 (5.)

- Using 600M parameter version without pre-trained weights improves MJI by 1.0 %

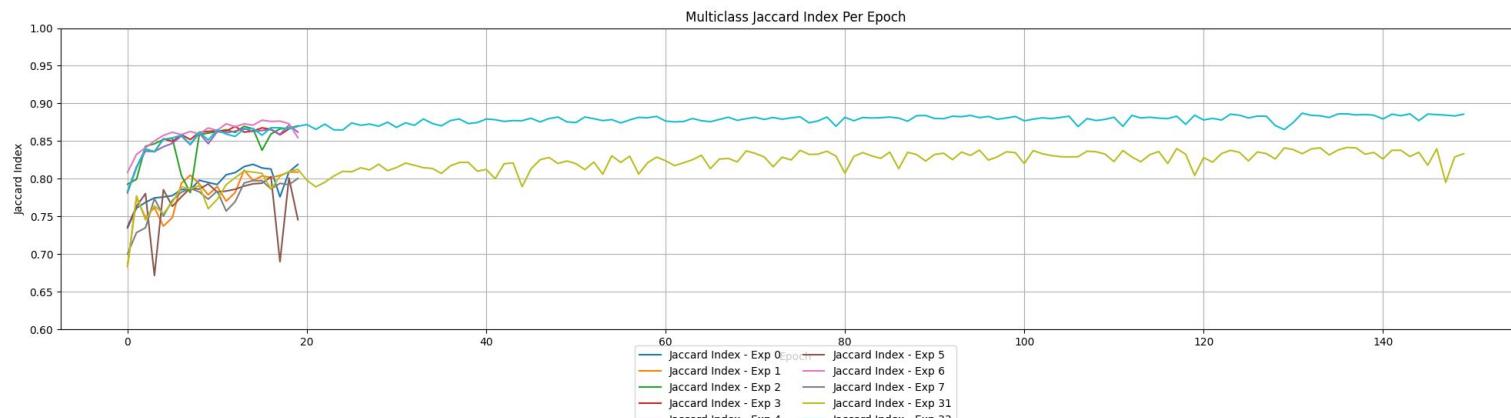
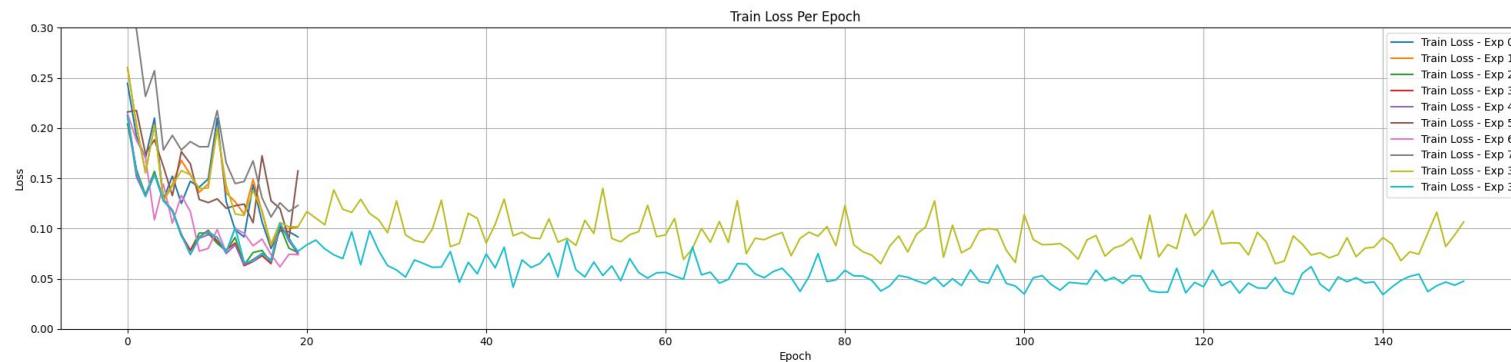
experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_3	v2_300	TRUE	[]	TRUE	FALSE	0.0901 (4.)	0.8576 (4.)
exp_6	v2_600	TRUE	[]	TRUE	FALSE	0. 0711 (1.)	0.8836 (1.)

- Using 600M parameter version with pre-trained weights improves MJI by 3.0 %

experiment	backbone	backbone_pret rained	backbone_enco ding	freeze_back bone	freeze_deco der	Test Loss (ranking)	Test Multiclass Jaccard Index (ranking)
exp_1	v2_300	FALSE	[]	TRUE	FALSE	0.1191 (7.)	0.8103 (7.)
exp_7	v2_300	FALSE	[]	TRUE	TRUE	0.1386 (8.)	0.8041 (8.)

Training ONLY the head decreases MJI by 0.7 %

experiment	epochs	backbone	backbone_pr etrained	backbone_enc oding	freeze_bac kbone	freeze_dec oder	Test Loss	Test Multiclass Jaccard Index
exp_1	25	v2_300	FALSE	[]	TRUE	FALSE	0.1191	0.8103
exp_31	150	v2_300	FALSE	[]	TRUE	FALSE	0.0899	0.8450
exp_3	25	v2_300	TRUE	[]	TRUE	FALSE	0.0901	0.8576
exp_33	150	v2_300	TRUE	[]	TRUE	FALSE	0.0633	0.8876



## Effect of longer training (more epochs):

Without pre-trained backbone (exp 1 vs 31):

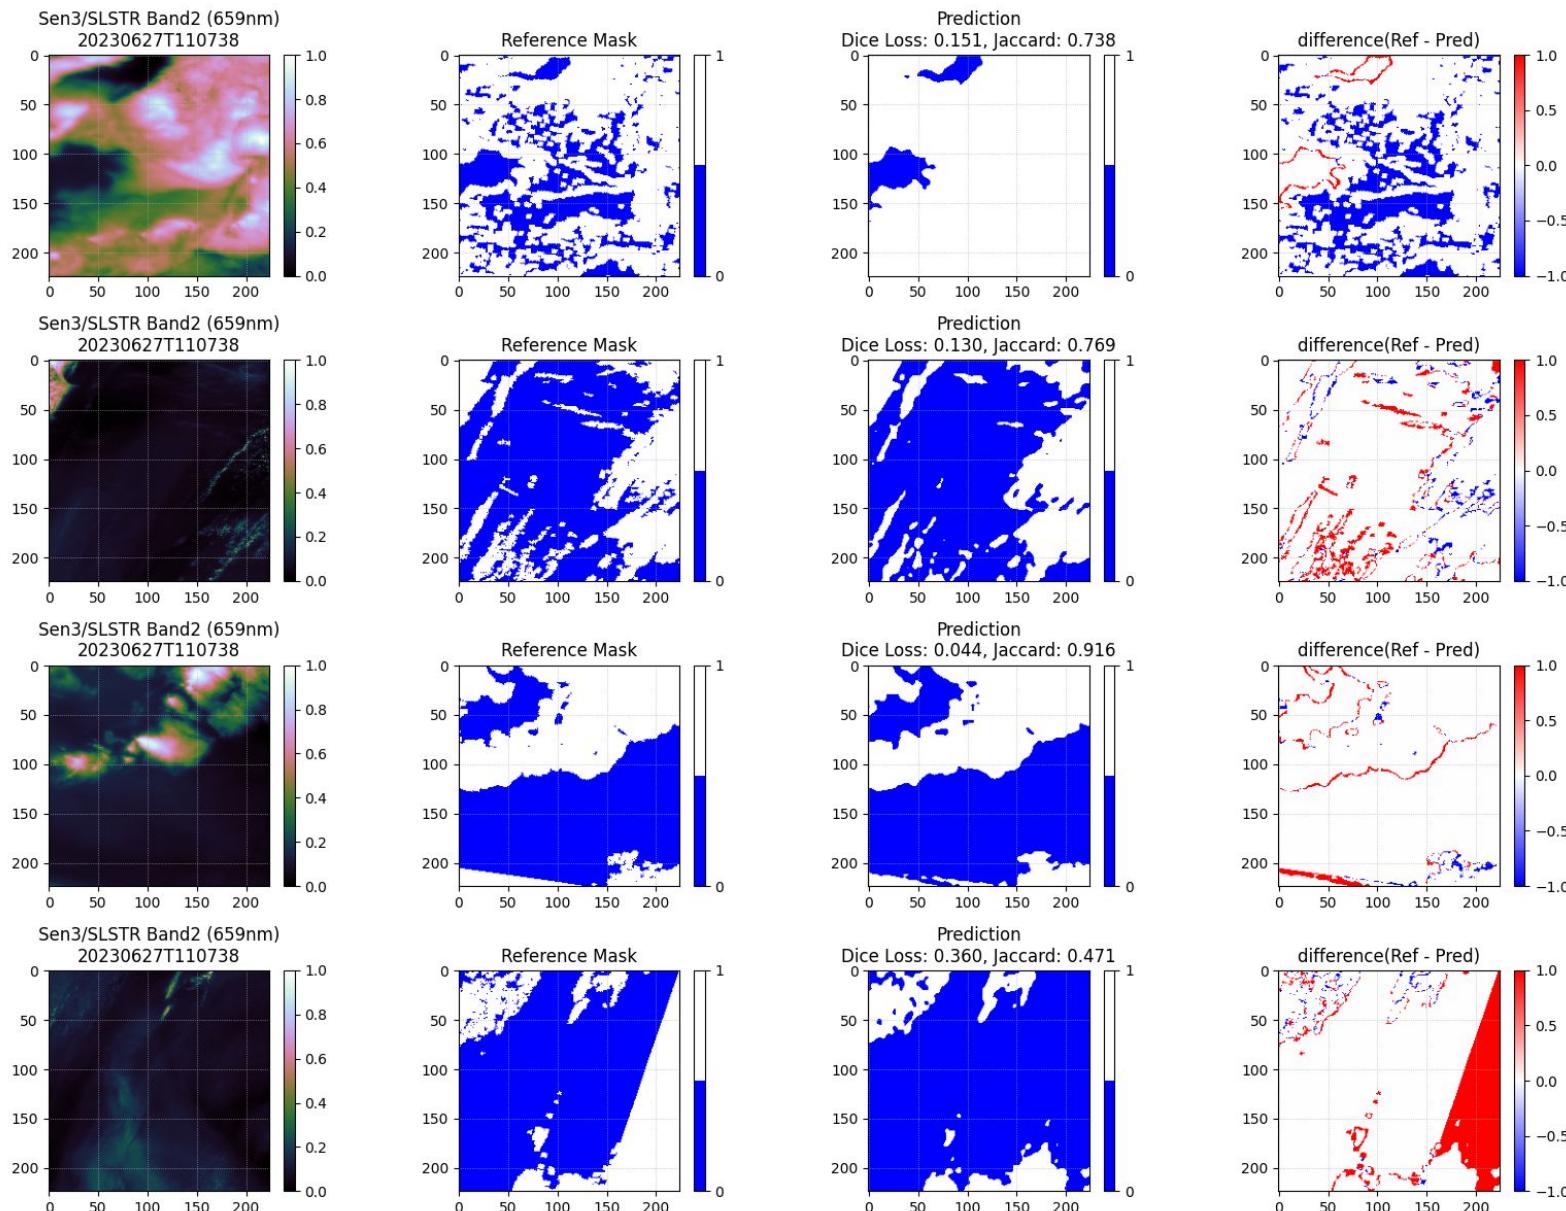
4 % improvement

BUT: still worse than with pre-trained

With pre-trained backbone (exp 3 vs 33):

2 % improvement

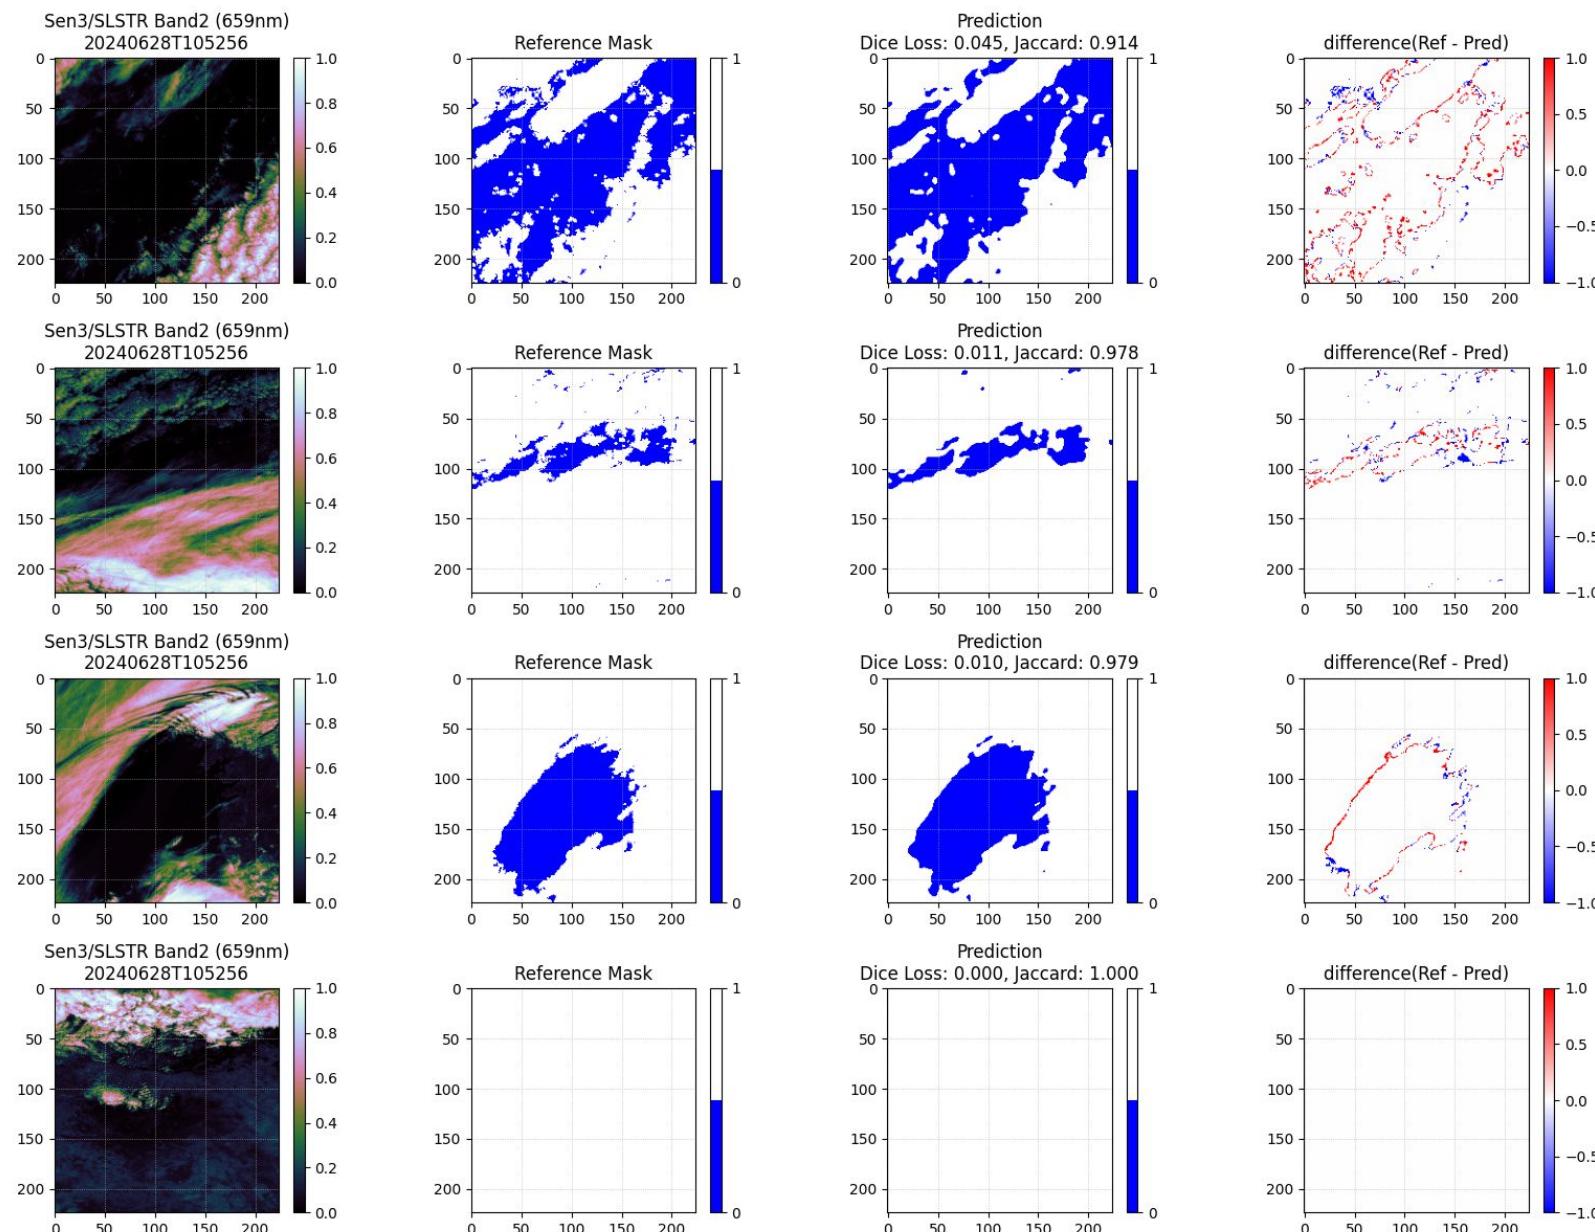
# Results: Visual examples from best model (6)



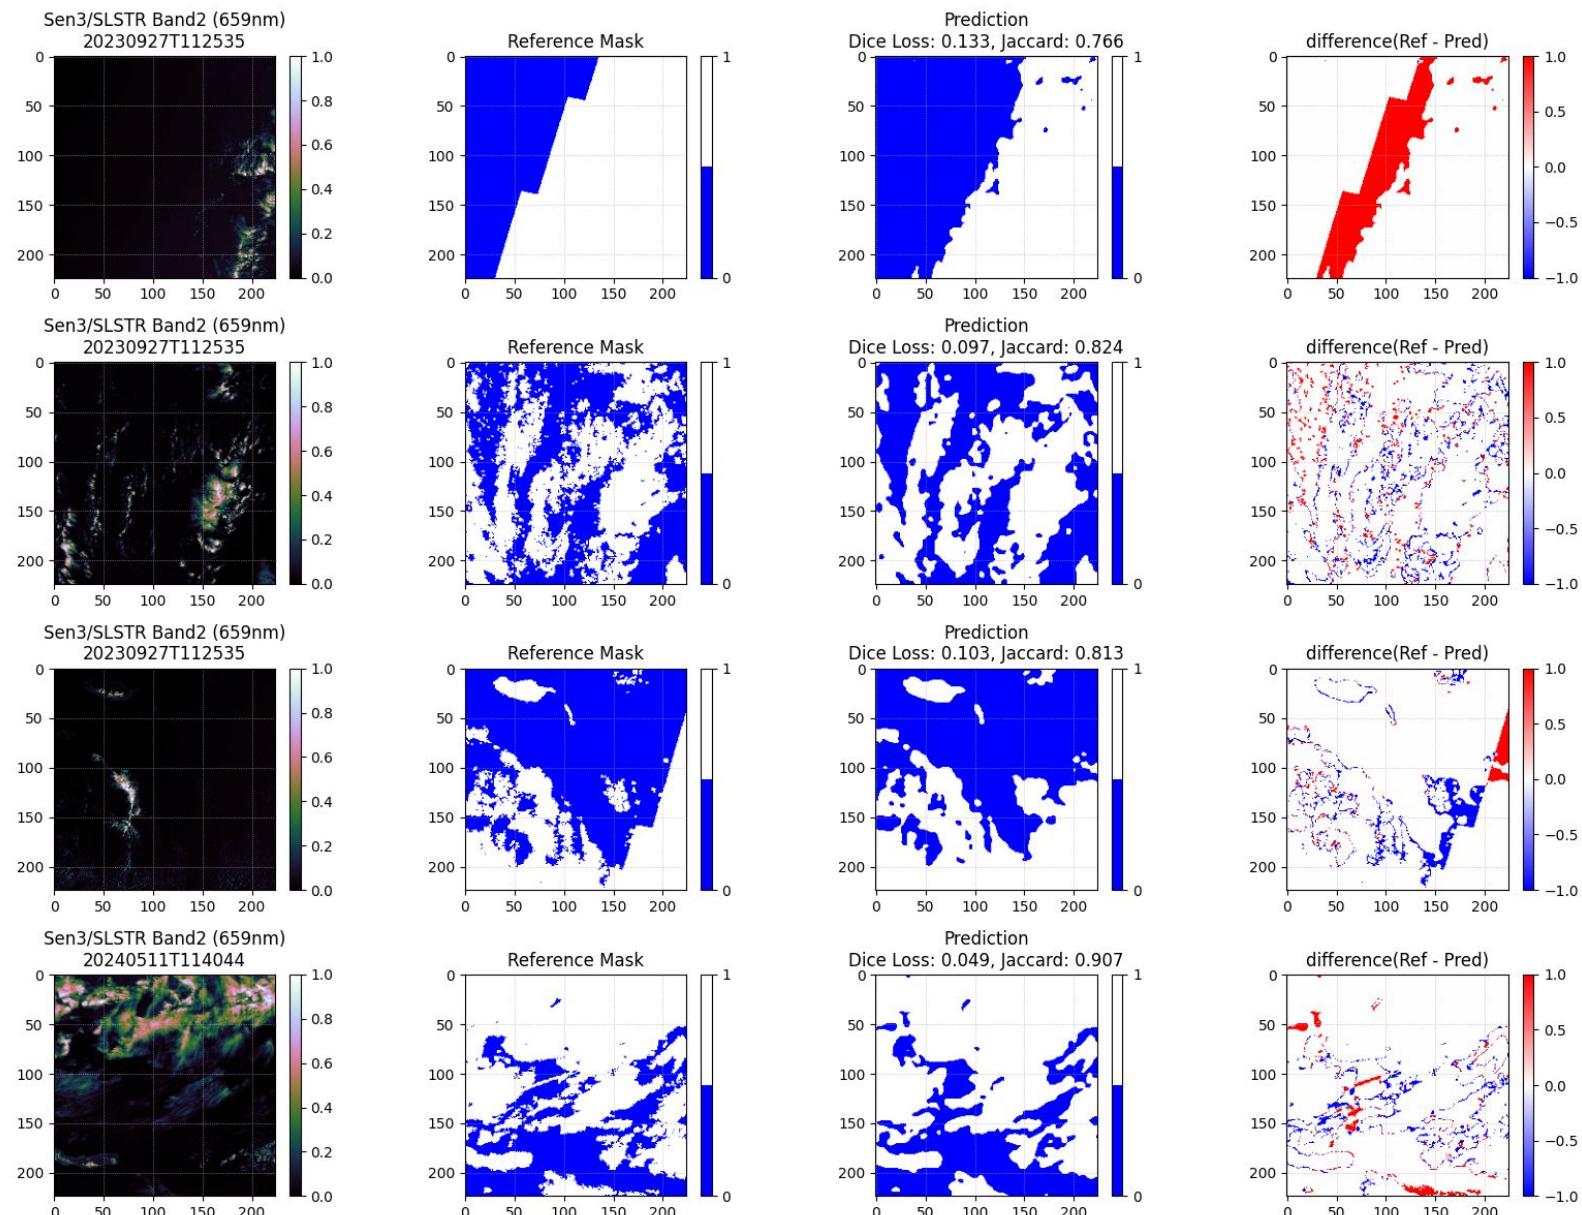


# Results: Visual examples from best model (6)

www.eumetsat.int



# Results: Visual examples from best model (6)



## Ablation Studies (exact improvements also depend on other settings):

- pre-trained weights, aka the foundation model itself improves accuracy ~ 5-6 %
- Training all weights improves accuracy ~1 - 2 %
- Using time and location embeddings improves accuracy ~ 2- 3 %
- Using more parameters improves MJI ~1 – 3 %
- Training only the head (without pre-trained weights) is only 9.8 % worse than the best model
- More epochs improve MJI ~2 – 4 %

## Workflow:

- Not only improvement of performance through pre-trained weights, but also python framework (terratorch) makes it easy and quick to use
- As always in ML: large part of the task is pre-processing the dataset

## Other advantages:

- Computation probably much faster, especially if a GPU is available

## Research Questions and Answers

- How does a foundation model perform that was pre-trained with a higher resolution?
  - It can be used and improves the accuracy!
- How does a foundation model perform that was pre-trained on other spectral bands?
  - It can be used and improves the accuracy, even when choosing suboptimal corresponding spectral bands!
- How easy is it to apply foundation models? What computational and intellectual resources are required?
  - One 20GB GPU is enough, if the individual training data points are small (here: 224x224x6).

## Pre-Processing/ Data curation:

- Maybe tested if the conversion Radiance to Reflectance is necessary for the model – it doesn't seem like it.
- Wrong corresponding SLSTR bands in dataset
- Can't cut out what I want/ unflexible cut out because of regridding
- Too much Land data involved ☐ very easy to predict

## Model experiments:

- Test with data from different tiles

## Evaluation:

- Predictions are overly smoothed ☐ common ML problem!
- No comparison to “traditional” ML
  - Only comparison between encoder with pre-trained weights and encoder with random weights, which basically is the comparison between using the foundation model or without, but not a comparison to a different baseline ML model



# Future steps

- Improve model accuracy even further by:
  - More data
  - Find 'best' configuration
- Do more ablation studies to understand the behaviour:
  - How does it perform with less data?
  - How does it perform on other locations and times?
  - How important is the class balance?
  - How flexible is the model with the used spectral bands?
  - Test it with other, even coarser instruments like SEVIRI (Meteosat)
- Adapt training data by hand and test against numerical algorithm for smoke cases



- Bommasani et al. 2022 – On the Opportunities and Risks of Foundation Models
- Jakubik et al. 2023 – Foundation Models for Generalist Geospatial Artificial Intelligence (Prithvi 1.0)
- Szwarcman et al. 2025 - Prithvi-EO-2.0: A Versatile Multi-Temporal Foundation Model for Earth Observation Applications
- Zhu et al. 2024 – On the Foundations of Earth and Climate Foundation Models



**Thank you!**

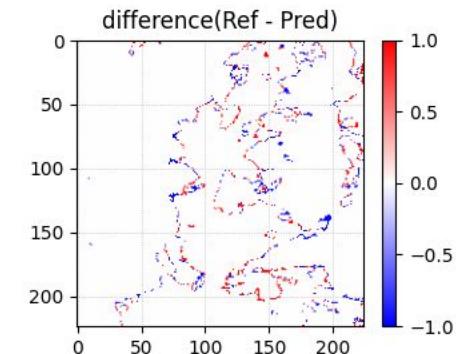
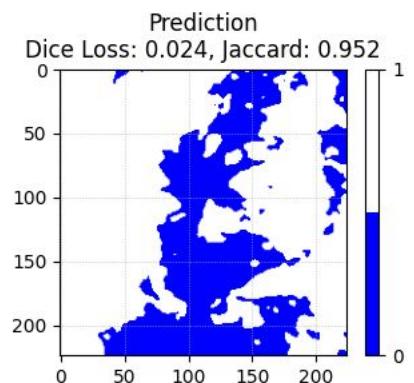
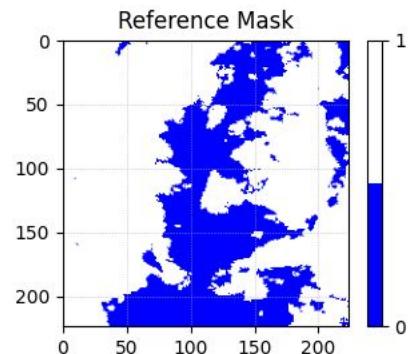
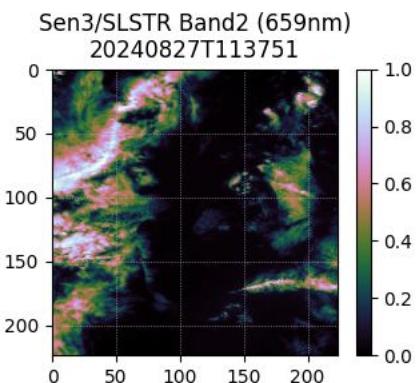
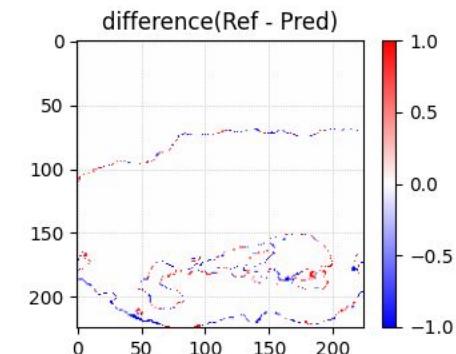
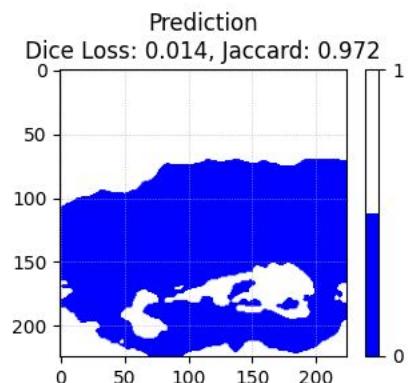
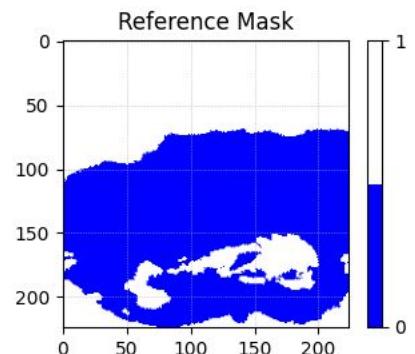
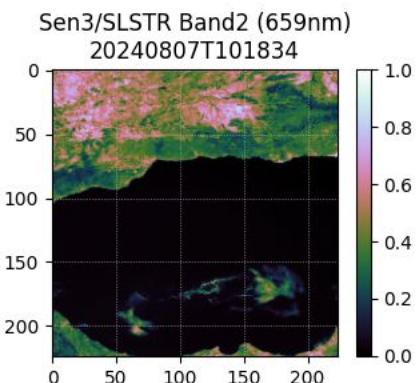
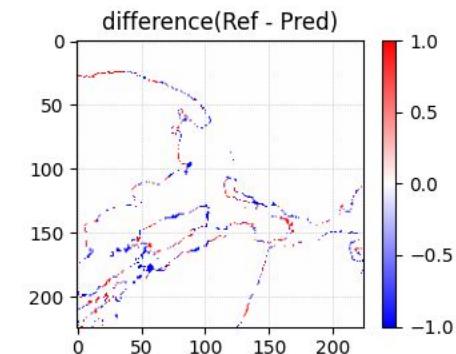
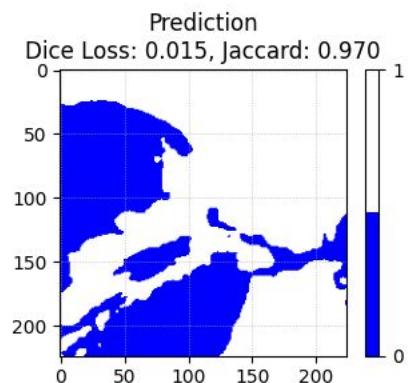
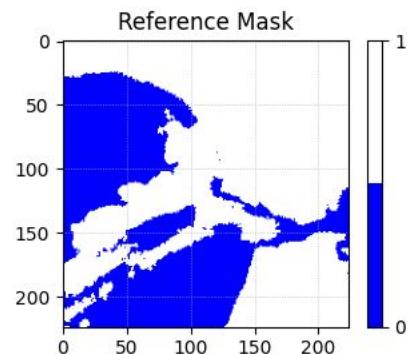
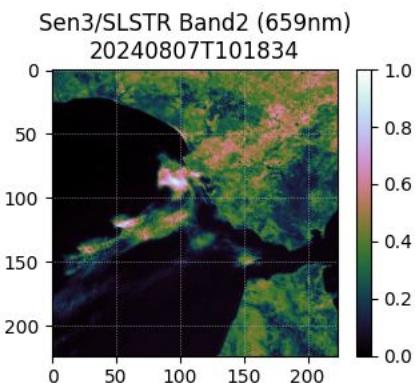
Questions are welcome.

Contact:

Boran Frank, Anna-Lena Erdmann (LinkedIn)

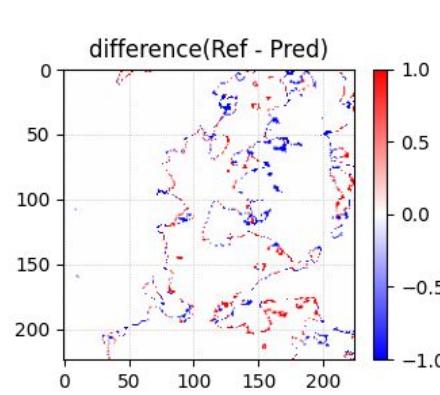
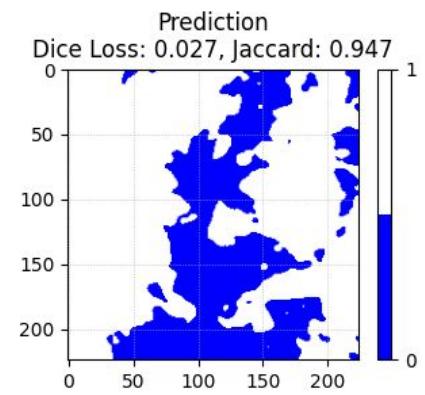
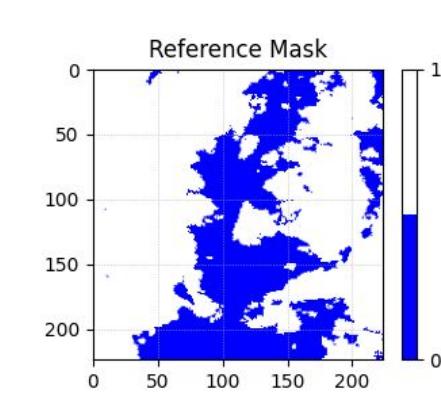
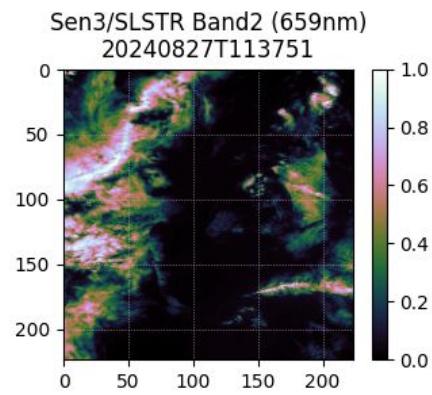
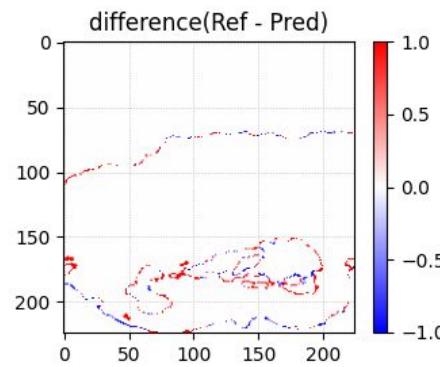
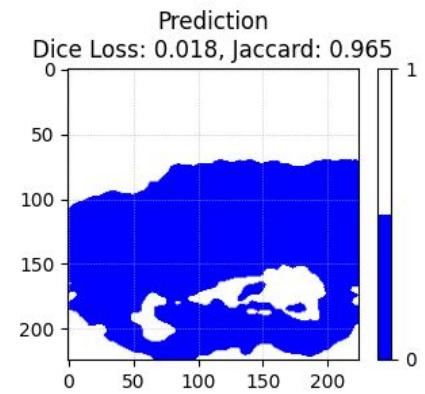
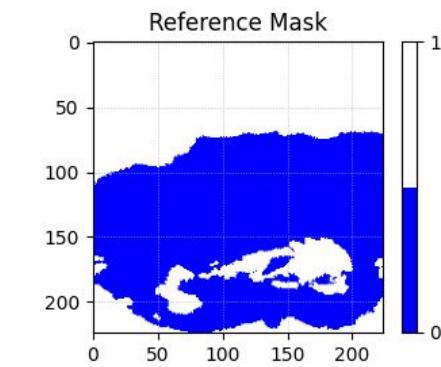
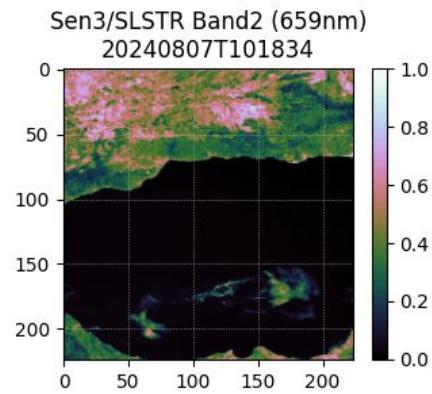
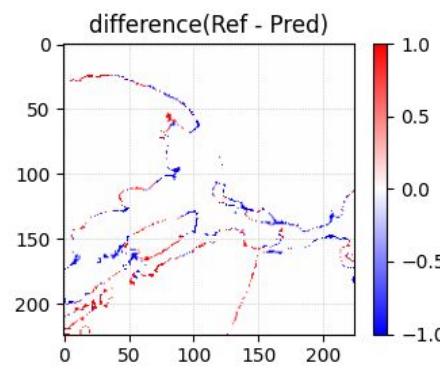
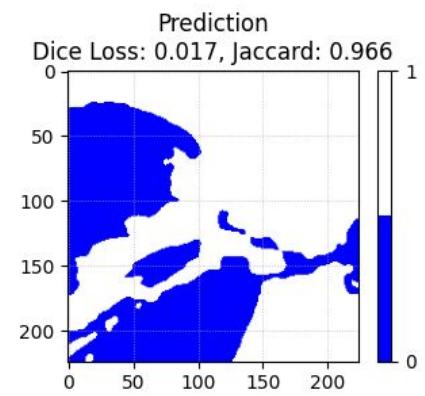
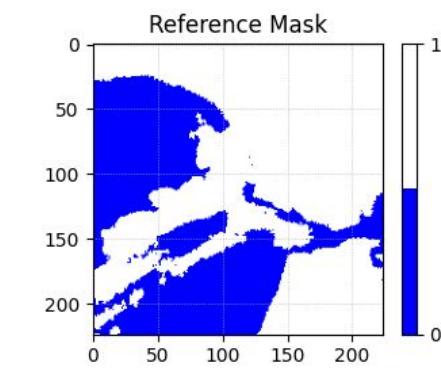
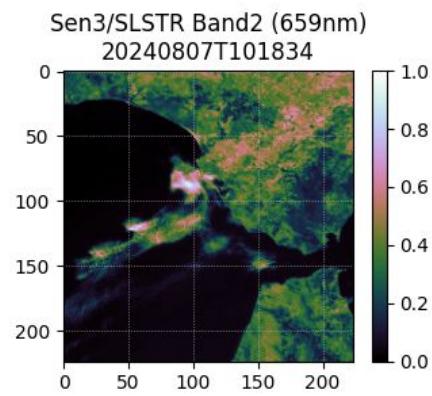
# Results: Visual comparisons

Exp\_33:



# Results: Visual comparisons

Exp\_6:



# Results: Visual comparisons

Exp\_7:  
Only  
head, no  
pre-traine  
d  
backbone

